

Sequencing Legal DNA

NLP for Law and Political Economy

5. Neural Nets and Word Embeddings

Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$, a vector of predicted probabilities across classes for each document i .

Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$, a vector of predicted probabilities across classes for each document i .
- ▶ This vector is a compressed representation of the outcome-predictive text features \mathbf{x}_i

Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$, a vector of predicted probabilities across classes for each document i .
- ▶ This vector is a compressed representation of the outcome-predictive text features \mathbf{x}_i
 - ▶ \mathbf{x}_i is itself a compressed representation of the unprocessed document \mathcal{D}_i .

Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$, a vector of predicted probabilities across classes for each document i .
- ▶ This vector is a compressed representation of the outcome-predictive text features \mathbf{x}_i
 - ▶ \mathbf{x}_i is itself a compressed representation of the unprocessed document \mathcal{D}_i .
- ▶ Correspondingly: the parameters $\hat{\theta}$ can also be understood as a compressed (or “learned”) representation:
 - ▶ it contains information about the training corpus, the text features, and the outcomes.

Information in $\hat{\theta}$

- ▶ Say we train a multinomial logistic regression on a bag-of-words representation of the documents.
- ▶ Let θ be the learned matrix of parameters relating words to outcomes:
 - ▶ It contains n_y columns = n_x -vectors representing the outcome classes.

Information in $\hat{\theta}$

- ▶ Say we train a multinomial logistic regression on a bag-of-words representation of the documents.
- ▶ Let θ be the learned matrix of parameters relating words to outcomes:
 - ▶ It contains n_y columns = n_x -vectors representing the outcome classes.
 - ▶ It contains n_x rows = n_y -vectors representing each word in the vocabulary.

Information in $\hat{\theta}$

- ▶ Say we train a multinomial logistic regression on a bag-of-words representation of the documents.
- ▶ Let θ be the learned matrix of parameters relating words to outcomes:
 - ▶ It contains n_y columns = n_x -vectors representing the outcome classes.
 - ▶ It contains n_x rows = n_y -vectors representing each word in the vocabulary.
- ▶ How to use this?
 - ▶ could cluster column vectors to understand which outcomes are similar/related.
 - ▶ could cluster row vectors to understand which features are similar/related.

Preview of Word Embeddings

- ▶ Let's say \mathbf{x}_i is a bag-of-words representation for document i with length n_i . We can write

$$\mathbf{x}_i = \frac{1}{n_i} \sum_{l=1}^{n_i} \mathbf{x}_i^{[l]}$$

- ▶ l indexes words in the the document
- ▶ each vector $\mathbf{x}_i^{[l]}$ is an n_x -dimensional one-hot vector – all entries are zero except the single entry corresponding to the word at l , which is 1.

Preview of Word Embeddings

- ▶ Let's say \mathbf{x}_i is a bag-of-words representation for document i with length n_i . We can write

$$\mathbf{x}_i = \frac{1}{n_i} \sum_{l=1}^{n_i} \mathbf{x}_i^{[l]}$$

- ▶ l indexes words in the the document
- ▶ each vector $\mathbf{x}_i^{[l]}$ is an n_x -dimensional one-hot vector – all entries are zero except the single entry corresponding to the word at l , which is 1.
- ▶ Now let $\theta^{[l]}$ be the row of θ corresponding to the word w_l . We can write

$$\hat{\mathbf{y}}_i = \frac{1}{n_i} \sum_{l=1}^{n_i} \theta^{[l]}$$

the sum of the n_y -dimensional word representations (the row vectors from above).

- ▶ this is called the “continuous bag of words (CBOW)” representation.
- ▶ θ is a word embedding matrix.

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings

Word2Vec and GloVe

“Neural Networks”

- ▶ “Neural”:
 - ▶ nothing like brains

“Neural Networks”

- ▶ “Neural”:
 - ▶ nothing like brains
- ▶ “Networks”:
 - ▶ nothing to do with “networks” as normally understood – in particular, nothing to do with network theory in math or social science.

Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.

Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.

Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.
- ▶ Training algorithms have improved – small tweaks have made a huge impact.

Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.
- ▶ Training algorithms have improved – small tweaks have made a huge impact.
- ▶ Some theoretical limitations of NNs have turned out to be benign in practice – for example, they work well on non-convex functions.

Will it last?

- ▶ Three key principles of deep learning will persist:
 - ▶ **Simplicity**
 - ▶ feature engineering is obsolete
 - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.

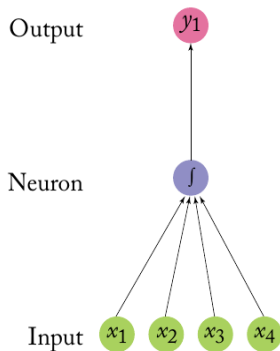
Will it last?

- ▶ Three key principles of deep learning will persist:
 - ▶ **Simplicity**
 - ▶ feature engineering is obsolete
 - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.
 - ▶ **Scalability**
 - ▶ amenable to parallelization on GPUs or TPUs (tensor processing units)
 - ▶ trained on batches of data, so can be scaled to datasets of arbitrary size.

Will it last?

- ▶ Three key principles of deep learning will persist:
 - ▶ **Simplicity**
 - ▶ feature engineering is obsolete
 - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.
 - ▶ **Scalability**
 - ▶ amenable to parallelization on GPUs or TPUs (tensor processing units)
 - ▶ trained on batches of data, so can be scaled to datasets of arbitrary size.
 - ▶ **Versatility and reusability**
 - ▶ can be trained on additional data without restarting from scratch, therefore amenable for continuous online learning.
 - ▶ deep-learning models are repurposable and thus reusable

A “Neuron”



- ▶ A neuron multiplies each input by its weight, sums them, applies a non-linear function to the result, and passes the output.
 - ▶ e.g., the \int shape indicates a sigmoid transformation.

In Notation

- ▶ The simplest neural network is called a perceptron:

$$\text{MLP0}(\mathbf{x}) = \mathbf{x} \cdot \boldsymbol{\omega}$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega} \in \mathbb{R}^{n_x \times n_y}$$

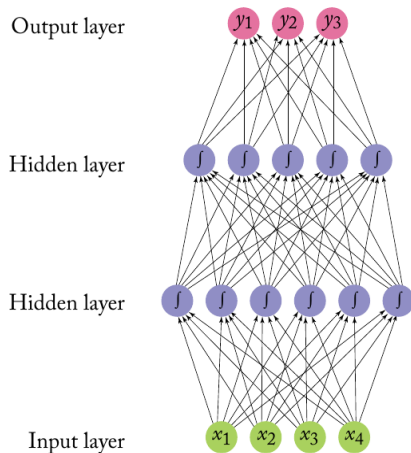
here, $\boldsymbol{\omega}$ is the matrix of weights in the layer.

- ▶ In more standard notation, there would be an additional constant (or “bias”) term:

$$\text{MLP0}(\mathbf{x}) = \alpha + \mathbf{x} \cdot \boldsymbol{\omega}$$

- ▶ We leave it out by assuming that \mathbf{x} is de-meant or has an extra column of ones.

A Feed-Forward Neural Network



- A feed-forward network is simply a stack of linear models, separated by non-linear functions.

Multi-Layer Perceptron

- ▶ An multi-layer perceptron (MLP) with one hidden layer is

$$\text{MLP1}(\mathbf{x}) = \mathbf{g}(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_y},$$

- ▶ n_1 = dimensionality in first (and only) hidden layer
- ▶ $\boldsymbol{\omega}_1$ = set of learnable weights for the first linear transformation of the inputs.
- ▶ $\mathbf{g}(\cdot)$ = an element-wise non-linear function (an “activation function”)
- ▶ $\boldsymbol{\omega}_2$ = weights on the second linear transformation leading to the output.

Multi-Layer Perceptron

- ▶ An multi-layer perceptron (MLP) with one hidden layer is

$$\text{MLP1}(\mathbf{x}) = \mathbf{g}(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_y},$$

- ▶ n_1 = dimensionality in first (and only) hidden layer
- ▶ $\boldsymbol{\omega}_1$ = set of learnable weights for the first linear transformation of the inputs.
- ▶ $\mathbf{g}(\cdot)$ = an element-wise non-linear function (an “activation function”)
- ▶ $\boldsymbol{\omega}_2$ = weights on the second linear transformation leading to the output.
- ▶ MLP1 can approximate any continuous function on a closed and bounded subset of \mathbb{R}^n , and any mapping from one finite discrete space to another finite discrete space (Hornik et al 1989, Cybenko 1989).
 - ▶ But MLP1 would have to be exponentially large in some cases (Telgarsky 2016) .

Two hidden layers

- ▶ Adding a second hidden layer gives

$$\text{MLP2}(\mathbf{x}) = \mathbf{g}_2(\mathbf{g}_1(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2) \cdot \boldsymbol{\omega}_3$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_2}, \boldsymbol{\omega}_3 \in \mathbb{R}^{n_2 \times n_y}$$

- ▶ n_2 = number of neurons in second hidden layer.

Two hidden layers

- ▶ Adding a second hidden layer gives

$$\text{MLP2}(\mathbf{x}) = \mathbf{g}_2(\mathbf{g}_1(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2) \cdot \boldsymbol{\omega}_3$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_2}, \boldsymbol{\omega}_3 \in \mathbb{R}^{n_2 \times n_y}$$

- ▶ n_2 = number of neurons in second hidden layer.
- ▶ MLP2 can be written in the following decomposed notation:

$$\text{MLP2}(\mathbf{x}) =$$

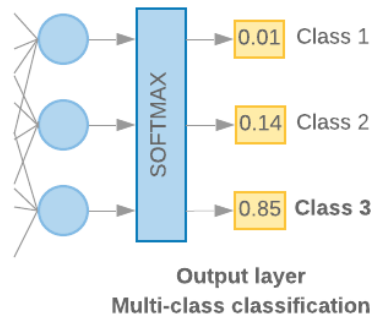

$$\mathbf{h}_1 = \mathbf{g}_1(\mathbf{x} \cdot \boldsymbol{\omega}_1)$$

$$\mathbf{h}_2 = \mathbf{g}_2(\mathbf{h}_1 \cdot \boldsymbol{\omega}_2)$$

$$\mathbf{y} = \mathbf{h}_2 \cdot \boldsymbol{\omega}_3$$

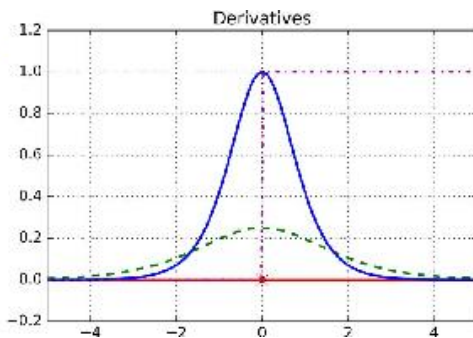
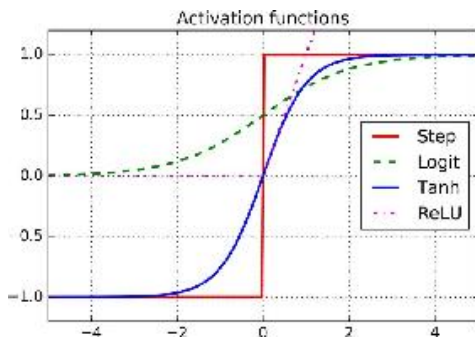
where \mathbf{h}_l give hidden layers.

Constructing the Last Layer



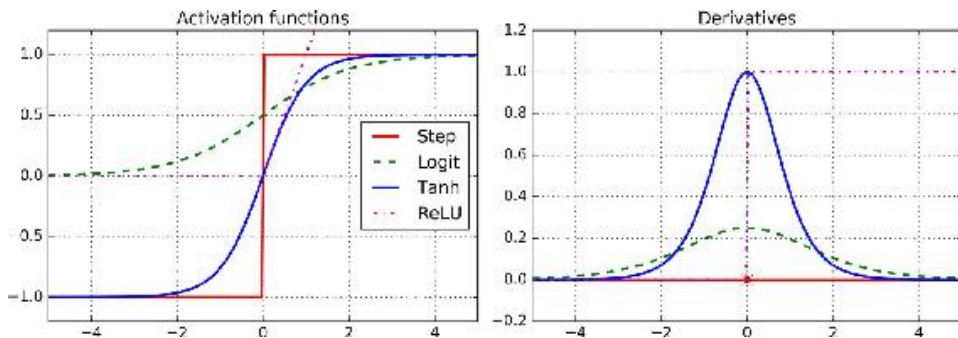
- ▶ MLPs will output a probability distribution across output classes.
 - ▶ can also output a real number, which would make a regression model.

What to pick for $g(\cdot)$



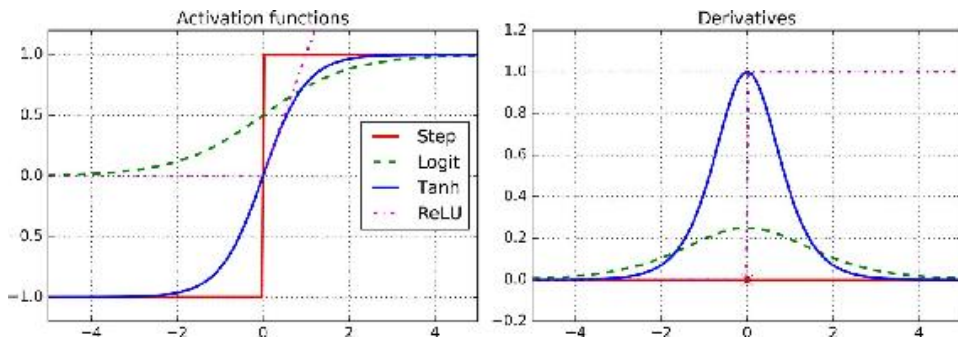
- logistic function: $\text{logit}(z) = \frac{1}{1+\exp(-z)}$

What to pick for $g(\cdot)$



- ▶ logistic function: $\text{logit}(z) = \frac{1}{1+\exp(-z)}$
- ▶ hyperbolic tangent function: $\tanh(z) = 2\sigma(2z) - 1$
 - ▶ ranges between -1 and 1 (rather than between 0 and 1, as the case with the logistic)
 - ▶ centered on zero, can speed up convergence

What to pick for $g(\cdot)$



- ▶ logistic function: $\text{logit}(z) = \frac{1}{1+\exp(-z)}$
- ▶ hyperbolic tangent function: $\tanh(z) = 2\sigma(2z) - 1$
 - ▶ ranges between -1 and 1 (rather than between 0 and 1, as the case with the logistic)
 - ▶ centered on zero, can speed up convergence
- ▶ ReLU (rectified linear unit) function: $\max\{0, z\}$,
 - ▶ deceptively simple, fast to compute, and very effective in practice
 - ▶ gradient does not saturate to zero for large values (but is flat below zero)

Backpropagation

- ▶ A crucial technology in deep learning is *backpropagation*, also known as *reverse-mode automatic differentiation (autodiff)*.
 - ▶ intuitively, computes the network's output error and how much each neuron contributes to the error, so they can be updated to reduce error.

Backpropagation

- ▶ A crucial technology in deep learning is *backpropagation*, also known as *reverse-mode automatic differentiation (autodiff)*.
 - ▶ intuitively, computes the network's output error and how much each neuron contributes to the error, so they can be updated to reduce error.
 - ▶ Keras/TensorFlow will do this under the hood, just like how scikit-learn does gradient descent.
 - ▶ Appendix D of the Geron book has a decent explanation for how it works.

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings

Word2Vec and GloVe

MLP baseline for Text Classification

Google Developers Advice

1. Calculate the number of samples/number of words per sample ratio.
2. If this ratio is less than 1500, tokenize the text as n-grams and use a simple multi-layer perceptron (MLP) model to classify them.
 - ▶ In the case of N-grams models, Google testers found that MLPs tended to out-perform logistic regression and gradient boosting machines.
 - ▶ A simple MLP is one of the models tried by Peterson and Spirling (2018).

Keras Basics

- ▶ See the Geron book and sample notebooks for Keras examples.
- ▶ “Dense” layer is the DNN baseline – means that all neurons are connected.

Keras Basics

- ▶ See the Geron book and sample notebooks for Keras examples.
- ▶ “Dense” layer is the DNN baseline – means that all neurons are connected.
- ▶ Output layer:
 - ▶ for regression, do not use an activation function
 - ▶ for binary classification, use `activation='sigmoid'`
 - ▶ for multi-class classification, use `activation='softmax'`

Loss function and metrics

- ▶ Loss function:
 - ▶ for regression, use `mean_squared_error`
 - ▶ for binary classification, use `binary_crossentropy`
 - ▶ for multi-class classification, use `sparse_categorical_crossentropy`

Loss function and metrics

- ▶ Loss function:
 - ▶ for regression, use `mean_squared_error`
 - ▶ for binary classification, use `binary_crossentropy`
 - ▶ for multi-class classification, use `sparse_categorical_crossentropy`
- ▶ Metrics:
 - ▶ for classification, can use accuracy and F_1
 - ▶ for regression, use R^2

Table 10-1. Typical regression MLP architecture

Hyperparameter	Typical value
# input neurons	One per input feature (e.g., $28 \times 28 = 784$ for MNIST)
# hidden layers	Depends on the problem, but typically 1 to 5
# neurons per hidden layer	Depends on the problem, but typically 10 to 100
# output neurons	1 per prediction dimension
Hidden activation	ReLU (or SELU, see Chapter 11)
Output activation	None, or ReLU/softplus (if positive outputs) or logistic/tanh (if bounded outputs)
Loss function	MSE or MAE/Huber (if outliers)

Table 10-2. Typical classification MLP architecture

Hyperparameter	Binary classification	Multilabel binary classification	Multiclass classification
Input and hidden layers	Same as regression	Same as regression	Same as regression
# output neurons	1	1 per label	1 per class
Output layer activation	Logistic	Logistic	Softmax
Loss function	Cross entropy	Cross entropy	Cross entropy

Initializing Neuron Weights

- ▶ Neuron weights have to be initialized randomly; otherwise they are collinear and backprop won't distinguish their contributions to the output error.

Initializing Neuron Weights

- ▶ Neuron weights have to be initialized randomly; otherwise they are collinear and backprop won't distinguish their contributions to the output error.
- ▶ Standard practice is to use “Glorot” or “He” initialization:

Activation function	Uniform distribution $[-r, r]$	Normal distribution
Logistic	$r = \sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
Hyperbolic tangent	$r = 4\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = 4\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
ReLU (and its variants)	$r = \sqrt{2}\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{2}\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$

Tuning NN Hyperparameters

- ▶ Number of neurons:
 - ▶ just pick 128 neurons per layer
 - ▶ overall, better to have too many neurons, and use regularization

Tuning NN Hyperparameters

- ▶ Number of neurons:
 - ▶ just pick 128 neurons per layer
 - ▶ overall, better to have too many neurons, and use regularization
- ▶ Number of hidden layers:
 - ▶ try between one and five.
 - ▶ adding layers usually helps more than adding neurons.

Tuning NN Hyperparameters

- ▶ Number of neurons:
 - ▶ just pick 128 neurons per layer
 - ▶ overall, better to have too many neurons, and use regularization
- ▶ Number of hidden layers:
 - ▶ try between one and five.
 - ▶ adding layers usually helps more than adding neurons.
- ▶ Activation functions:
 - ▶ ReLU is a good baseline in the hidden layers.

Tuning NN Hyperparameters

- ▶ Number of neurons:
 - ▶ just pick 128 neurons per layer
 - ▶ overall, better to have too many neurons, and use regularization
- ▶ Number of hidden layers:
 - ▶ try between one and five.
 - ▶ adding layers usually helps more than adding neurons.
- ▶ Activation functions:
 - ▶ ReLU is a good baseline in the hidden layers.
- ▶ See Geron, pp. 322-323 for tools to help with tuning, e.g. Keras Tuner.
 - ▶ See also Smith (2018).

Early stopping

An efficient training approach, that also works to regularize a model, is **early stopping**:

- ▶ Split data into three sets: training, validation, and test.

Early stopping

An efficient training approach, that also works to regularize a model, is **early stopping**:

- ▶ Split data into three sets: training, validation, and test.
- ▶ continually evaluate your model in validation set at regular intervals
- ▶ stop training when the validation-set accuracy starts to decrease.

Early stopping

An efficient training approach, that also works to regularize a model, is **early stopping**:

- ▶ Split data into three sets: training, validation, and test.
- ▶ continually evaluate your model in validation set at regular intervals
- ▶ stop training when the validation-set accuracy starts to decrease.
- ▶ evaluate model in test set.

Batch normalization

- ▶ Another trick to increase performance, speed up training, and regularize the model:
 - ▶ in between layers, zero-center and normalize the inputs to variance one.
 - ▶ normally done before a non-linear activation function

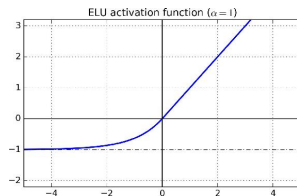
Batch normalization

- ▶ Another trick to increase performance, speed up training, and regularize the model:
 - ▶ in between layers, zero-center and normalize the inputs to variance one.
 - ▶ normally done before a non-linear activation function
- ▶ Does not work well in recurrent neural nets – can use gradient clipping (hard constraint on gradient) instead.

ELU and SELU

Exponential linear unit

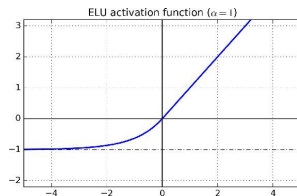
$$\text{ELU}(z) = \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$



ELU and SELU

Exponential linear unit

$$\text{ELU}(z) = \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$



SELU (Scaled ELU):

$$\text{SELU}(z) = \lambda \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$

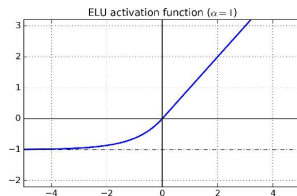
- ▶ Requires standardized inputs.
- ▶ Gaussian initialization of neuron weights:
 - ▶ mean zero and standard deviation = $n_x^{-1/2}$
- ▶ Set $\lambda \approx 1.0507, \alpha \approx 1.6732$.

- ▶ → Then layers will be self-normalizing; weights will converge to mean zero and variance once (Klambauer et al 2017)

ELU and SELU

Exponential linear unit

$$\text{ELU}(z) = \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$



SELU (Scaled ELU):

$$\text{SELU}(z) = \lambda \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$

- ▶ Requires standardized inputs.
- ▶ Gaussian initialization of neuron weights:
 - ▶ mean zero and standard deviation = $n_x^{-1/2}$
- ▶ Set $\lambda \approx 1.0507, \alpha \approx 1.6732$.

- ▶ → Then layers will be self-normalizing; weights will converge to mean zero and variance once (Klambauer et al 2017)

- ▶ In general, $\text{SELU} > \text{ELU} > \text{ReLU}$.

Optimizer and Learning Rate

- ▶ Choice of optimization algorithm is the topic of active research, which has shown that it can have a big impact on model performance.
 - ▶ See Geron, pp. 351-358. He recommends SGD with momentum, RMSProp, or Nadam.

Optimizer and Learning Rate

- ▶ Choice of optimization algorithm is the topic of active research, which has shown that it can have a big impact on model performance.
 - ▶ See Geron, pp. 351-358. He recommends SGD with momentum, RMSProp, or Nadam.
- ▶ On the learning rate, see Geron pp. 359-364, recommending 1cycle scheduling:
 - ▶ start at η_0 , increase linearly up to η_1 halfway through training, then decrease linearly to η_0 at the end.

Regularization for Sparse Models

- ▶ As with linear models, neural network parameters can be regularized with an L1 and/or L2 penalty to push weak neurons to zero and produce a sparse model.

Regularization for Sparse Models

- ▶ As with linear models, neural network parameters can be regularized with an L1 and/or L2 penalty to push weak neurons to zero and produce a sparse model.
- ▶ Another regularizer: “max-norm regularization”:
 - ▶ constrain each neuron’s weights ω by $\|\omega\|_2 \leq r$, where r is a hyperparameter.

Regularization for Sparse Models

- ▶ As with linear models, neural network parameters can be regularized with an L1 and/or L2 penalty to push weak neurons to zero and produce a sparse model.
- ▶ Another regularizer: “max-norm regularization”:
 - ▶ constrain each neuron's weights ω by $\|\omega\|_2 \leq r$, where r is a hyperparameter.
- ▶ But usually its better/simpler to use dropout.

Dropout

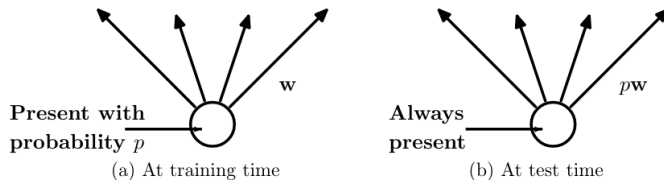


Figure 2: **Left:** A unit at training time that is present with probability p and is connected to units in the next layer with weights \mathbf{w} . **Right:** At test time, the unit is always present and the weights are multiplied by p . The output at test time is same as the expected output at training time.

Source: Srivastava et al, JMLR 2014

An elegant regularization technique:

- ▶ at every training step, every neuron has some probability (typically $p = 0.5$) of being temporarily dropped out, so that it will be ignored at this step.
- ▶ at test time, neurons don't get dropped anymore but coefficients are down-weighted by p .

Dropout

- ▶ Approximately equivalent to averaging the output of N models (where N is the number of neurons).

Dropout

- ▶ Approximately equivalent to averaging the output of N models (where N is the number of neurons).
- ▶ Neurons trained with dropout:
 - ▶ cannot co-adapt with neighboring neurons and must be independently useful.
 - ▶ cannot rely excessively on just a few input neurons; they have to pay attention to all input neurons.
 - ▶ makes the model less sensitive to slight changes in the inputs.

Dropout

- ▶ Approximately equivalent to averaging the output of N models (where N is the number of neurons).
- ▶ Neurons trained with dropout:
 - ▶ cannot co-adapt with neighboring neurons and must be independently useful.
 - ▶ cannot rely excessively on just a few input neurons; they have to pay attention to all input neurons.
 - ▶ makes the model less sensitive to slight changes in the inputs.
- ▶ If using SELU activation functions, use alpha dropout (Klambauer et al 2017).

Monte Carlo Dropout

- ▶ Normal dropout:
 - ▶ at test time, neurons don't get dropped but coefficients are down-weighted by p .

Monte Carlo Dropout

- ▶ Normal dropout:
 - ▶ at test time, neurons don't get dropped but coefficients are down-weighted by p .
- ▶ Monte Carlo dropout:
 - ▶ at test time, continue to allow dropout but produce 100 predictions, and average them.

Summary and Practical Guidelines

Table 11-3. Default DNN configuration

Hyperparameter	Default value
Kernel initializer	He initialization
Activation function	ELU
Normalization	None if shallow; Batch Norm if deep
Regularization	Early stopping (+ ℓ_2 reg. if needed)
Optimizer	Momentum optimization (or RMSProp or Nadam)
Learning rate schedule	1cycle

Table 11-4. DNN configuration for a self-normalizing net

Hyperparameter	Default value
Kernel initializer	LeCun initialization
Activation function	SELU
Normalization	None (self-normalization)
Regularization	Alpha dropout if needed
Optimizer	Momentum optimization (or RMSProp or Nadam)
Learning rate schedule	1cycle

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings

Word2Vec and GloVe

Application: Facts vs. Law

- ▶ A crucial distinction in law practice, and in applying caselaw, is to distinguish the law from the facts.
- ▶ In Cao, Ash, Chen (2018), we made a sample of fact and law sections in court cases and trained a model to predict them based on the text.

Application: Facts vs. Law

- ▶ A crucial distinction in law practice, and in applying caselaw, is to distinguish the law from the facts.
- ▶ In Cao, Ash, Chen (2018), we made a sample of fact and law sections in court cases and trained a model to predict them based on the text.
- ▶ Corpus:
 - ▶ all cases from courtlistener.com with an annotated “FACTS” section (extracted using regular expressions).
 - ▶ 23,497 sections, split into 1.3M paragraphs
 - ▶ 36% facts, 64% laws

Models and Results

DEPN: Document represented as a sequence of word-dependency pairs, fed to a **MLP classifier** with two 500-dim hidden layers.

DEPW: a variant of DEPN using Shulayeva et al.'s (2017) dependency features.

SMITH : Implementation of Smith's (2014) method: Bag of Words, but filter words that are statistically related to either fact statements or law statements. Fed to a Logistic Regression classifier.

WS: Laver et al.'s (2003) Wordscore algorithm (cf. Sarel and Demirtas, 2017). Bag of Words, but each word assigned a score based on relation to fact or law. Document is represented as a (re-scaled) score that sums up the pre- calculated scores of the words it contains, weighted by their frequencies.

BOW : Bag of words representation of documents, fed to a Logistic Regression classifier.

D2V: 500-dim Doc2Vec vectors (Le and Mikolov, 2014), fed to **MLP**.

Models and Results

DEPN: Document represented as a sequence of word-dependency pairs, fed to a **MLP classifier** with two 500-dim hidden layers.

DEPW: a variant of DEPN using Shulayeva et al.'s (2017) dependency features.

SMITH : Implementation of Smith's (2014) method: Bag of Words, but filter words that are statistically related to either fact statements or law statements. Fed to a Logistic Regression classifier.

WS: Laver et al.'s (2003) Wordscore algorithm (cf. Sarel and Demirtas, 2017). Bag of Words, but each word assigned a score based on relation to fact or law. Document is represented as a (re-scaled) score that sums up the pre- calculated scores of the words it contains, weighted by their frequencies.

BOW : Bag of words representation of documents, fed to a Logistic Regression classifier.

D2V: 500-dim Doc2Vec vectors (Le and Mikolov, 2014), fed to **MLP**.

	F1 for facts	F1 for values
DEPN	73.38	74.66
DEPW	72.77	73.79
SMITH	71.85	71.4
WS	77.11	77.67
BOW	72.57	73.29
D2V	67.18	65.36

Table 2: 5-fold cross validation.

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings


Word2Vec and GloVe

Efficient Data Representations

Which of the following number sequences do you find the easiest to memorize?

- 40, 27, 25, 36, 81, 57, 10, 73, 19, 68
- 50, 48, 46, 44, 42, 40, 38, 36, 34, 32, 30, 28, 26, 24, 22, 20, 18, 16, 14

Autoencoders \leftrightarrow Optimal Compression Algorithms



- ▶ “Autoencoder” refers to a class of deep neural network that performs domain-specific compression and dimension reduction.
 - ▶ They learn efficient encodings of the data, which can then be decoded back to a *reconstruction* – a (minimally) lossy representation of the original data.
 - ▶ Can also randomly generate new data that look like the training data.

PCA is a linear autoencoder

```
from tensorflow import keras

encoder = keras.models.Sequential([keras.layers.Dense(2, input_shape=[3])])
decoder = keras.models.Sequential([keras.layers.Dense(3, input_shape=[2])])
autoencoder = keras.models.Sequential([encoder, decoder])

autoencoder.compile(loss="mse", optimizer=keras.optimizers.SGD(lr=0.1))
```

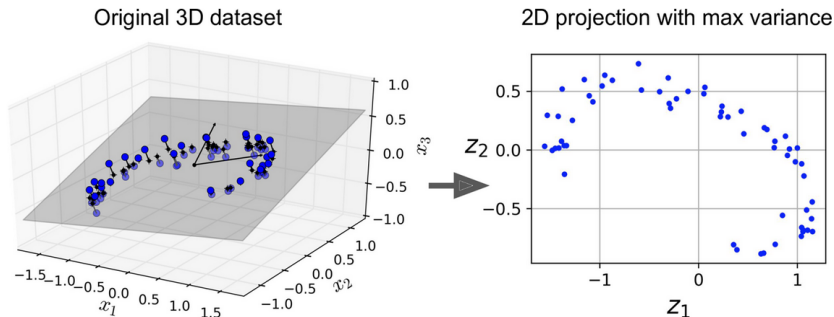


Figure 17-2. PCA performed by an undercomplete linear autoencoder




Figure 17-3. Stacked autoencoder

- ▶ Autoencoders work by stacking layers that gradually decrease in dimensionality to create the compressed representation (Z), and then gradually increase in dimensionality to try to reconstruct the input.
 - ▶ the autoencoder is implicitly solving the problem of maximizing entropy in the bottleneck layer.




Figure 17-3. Stacked autoencoder

- ▶ Autoencoders work by stacking layers that gradually decrease in dimensionality to create the compressed representation (Z), and then gradually increase in dimensionality to try to reconstruct the input.
 - ▶ the autoencoder is implicitly solving the problem of maximizing entropy in the bottleneck layer.
 - ▶ for symmetric autoencoders, **tying weights** of the encoding and decoding segments will speed up training and tends to improve performance.

Reconstruction from encoded vector



Figure 17-4. Original images (top) and their reconstructions (bottom)

Autoencoders for Data Visualization




Figure 17-5. Fashion MNIST visualization using an autoencoder followed by t-SNE

- ▶ Decent baseline for visualizing the encodings:
 - ▶ use an autoencoder to compress your data to relatively low dimension (e.g. 32 dimensions)
 - ▶ then use t-SNE for mapping the compressed data to a 2D plane.

Denoising Autoencoders




Figure 17-8. Denoising autoencoders, with Gaussian noise (left) or dropout (right)

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings

Word2Vec and GloVe

What is an Embedding?

- ▶ A (relatively) low-dimensional vector representation of a categorical variable.
 - ▶ spatial location of the vector encodes predictive information about the category.

What is an Embedding?

- ▶ A (relatively) low-dimensional vector representation of a categorical variable.
 - ▶ spatial location of the vector encodes predictive information about the category.
- ▶ e.g., trying to predict how employment responds to economic growth with data from U.S. states:
 - ▶ instead of including a fifty-dimensional categorical variable, include two-dimensional latitude and longitude

What is an Embedding?

- ▶ A (relatively) low-dimensional vector representation of a categorical variable.
 - ▶ spatial location of the vector encodes predictive information about the category.
- ▶ e.g., trying to predict how employment responds to economic growth with data from U.S. states:
 - ▶ instead of including a fifty-dimensional categorical variable, include two-dimensional latitude and longitude
 - ▶ or initialize each state to a random two-dimensional vector, and let the model decide where to move the states to improve prediction on your task (e.g.).

An embedding layer is just matrix multiplication

- ▶ An embedding layer can be represented as

$$x = v(w)$$
$$\underbrace{x}_{n_E \times 1} = \underbrace{E}_{n_E \times n_w} \cdot \underbrace{w}_{n_w \times 1}$$

- ▶ w , a categorical variable (e.g., representing a word)
 - ▶ one-hot vector with a single item equaling one.
 - ▶ The input to the embedding layer.

An embedding layer is just matrix multiplication

- ▶ An embedding layer can be represented as

$$x = v(w)$$
$$\underbrace{x}_{n_E \times 1} = \underbrace{E}_{n_E \times n_w} \cdot \underbrace{w}_{n_w \times 1}$$

- ▶ w , a categorical variable (e.g., representing a word)
 - ▶ one-hot vector with a single item equaling one.
 - ▶ The input to the embedding layer.
- ▶ x , a dense representation of the variable.
 - ▶ The output of the embedding layer.

An embedding layer is just matrix multiplication

- ▶ An embedding layer can be represented as

$$x = v(w)$$
$$\underbrace{x}_{n_E \times 1} = \underbrace{E}_{n_E \times n_w} \cdot \underbrace{w}_{n_w \times 1}$$

- ▶ w , a categorical variable (e.g., representing a word)
 - ▶ one-hot vector with a single item equaling one.
 - ▶ The input to the embedding layer.
- ▶ x , a dense representation of the variable.
 - ▶ The output of the embedding layer.
- ▶ An embedding function $v(\cdot)$, defined by matrix E , learnable by the DNN

The Embedding Matrix E

- ▶ The model learns the weights of the embedding matrix in the same way that it would learn any model parameters.
- ▶ The rows of the matrix correspond to vectors for the n_w categories.
 - ▶ These are the “word vectors” that people talk about when they mention word embeddings or Word2Vec.







Figure 13-4. Embeddings will gradually improve during training







Embedding Layers versus Dense Layers

- ▶ An embedding layer is statistically equivalent to a fully-connected dense layer with sparse data set as input and linear activation.

Embedding Layers versus Dense Layers

- ▶ An embedding layer is statistically equivalent to a fully-connected dense layer with sparse data set as input and linear activation.
- ▶ Why use an embedding layer rather than a dense layer?
 - ▶ embedding layers are much faster for this purpose
 - ▶ batch updating with regularization and dropout do not work well on sparse data.

Embedding Layers versus Dense Layers

- ▶ An embedding layer is statistically equivalent to a fully-connected dense layer with sparse data set as input and linear activation.
- ▶ Why use an embedding layer rather than a dense layer?
 - ▶ embedding layers are much faster for this purpose
 - ▶ batch updating with regularization and dropout do not work well on sparse data.
- ▶ Geron (pg. 433) advises using one-hot encoding for less than 10 items, embeddings for more than 50, and in between it's unclear.

MLP output layer as embedding matrix

- ▶ Consider the output layer for an MLP:
 - ▶ last hidden layer l , with dimensionality n_l
 - ▶ output classes with dimensionality n_y
 - ▶ layer $l = n_l \times n_y$ matrix of weights E_y

MLP output layer as embedding matrix

- ▶ Consider the output layer for an MLP:
 - ▶ last hidden layer l , with dimensionality n_l
 - ▶ output classes with dimensionality n_y
 - ▶ layer $l = n_l \times n_y$ matrix of weights E_y
- ▶ columns of E_y give n_y outcome class vectors.
 - ▶ vector similarities between columns indicate model's learned similarities between the output classes (Goldberg pg. 99)

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings

Word2Vec and GloVe

Word Embeddings

- ▶ Word embeddings are NN layers that map word indexes to dense vectors.

Word Embeddings

- ▶ Word embeddings are NN layers that map word indexes to dense vectors.
- ▶ Documents are lists of word indexes $\{w_1, w_2, \dots, w_{n_i}\}$.
 - ▶ equivalently, let w_i be a one-hot vector (dimensionality $n_w = \text{vocab size}$) where the associate word's index equals one.

Word Embeddings

- ▶ Word embeddings are NN layers that map word indexes to dense vectors.
- ▶ Documents are lists of word indexes $\{w_1, w_2, \dots, w_{n_i}\}$.
 - ▶ equivalently, let w_i be a one-hot vector (dimensionality $n_w = \text{vocab size}$) where the associate word's index equals one.
 - ▶ Normalize all documents to the same length L ; shorter documents can be padded with a null token. This requirement can be relaxed with recurrent neural networks.

Word Embeddings

- ▶ Word embeddings are NN layers that map word indexes to dense vectors.
- ▶ Documents are lists of word indexes $\{w_1, w_2, \dots, w_{n_i}\}$.
 - ▶ equivalently, let w_i be a one-hot vector (dimensionality $n_w = \text{vocab size}$) where the associate word's index equals one.
 - ▶ Normalize all documents to the same length L ; shorter documents can be padded with a null token. This requirement can be relaxed with recurrent neural networks.
- ▶ The embedding layer replaces the list of sparse one-hot vectors with a list of n_E -dimensional ($n_E \ll n_w$) dense vectors

$$\mathbf{X} = \begin{bmatrix} x_1 & \dots & x_L \end{bmatrix}$$

where

$$\underbrace{x_j}_{n_E \times 1} = \underbrace{E}_{n_E \times n_w} \cdot \underbrace{w_j}_{n_w \times 1}$$

Word Embeddings

- ▶ Word embeddings are NN layers that map word indexes to dense vectors.
- ▶ Documents are lists of word indexes $\{w_1, w_2, \dots, w_{n_i}\}$.
 - ▶ equivalently, let w_i be a one-hot vector (dimensionality $n_w = \text{vocab size}$) where the associate word's index equals one.
 - ▶ Normalize all documents to the same length L ; shorter documents can be padded with a null token. This requirement can be relaxed with recurrent neural networks.
- ▶ The embedding layer replaces the list of sparse one-hot vectors with a list of n_E -dimensional ($n_E \ll n_w$) dense vectors

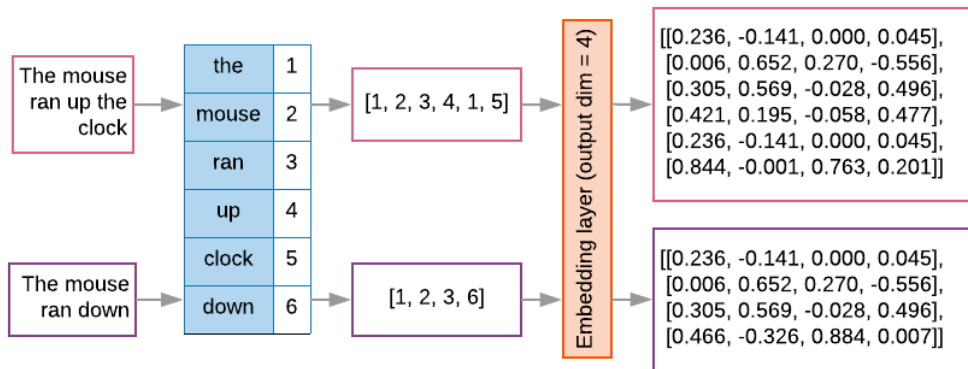
$$\mathbf{X} = \begin{bmatrix} x_1 & \dots & x_L \end{bmatrix}$$

where

$$\underbrace{x_j}_{n_E \times 1} = \underbrace{E}_{n_E \times n_w} \cdot \underbrace{w_j}_{n_w \times 1}$$

- ▶ This \mathbf{X} matrix is then flattened into an $L * n_E$ vector for input to the next layer.

Illustration



Continuous Bag-of-Words Representation

- ▶ Let w_j be the one-hot representation of feature j
- ▶ The dot product $w_j \cdot \mathbf{E}$ selects the row of \mathbf{E} corresponding to j :

$$v(w_j) = \mathbf{w}_j \cdot \mathbf{E}$$

- ▶ And a document can be represented as

$$\mathbf{x}_i = \text{CBOW}(x_1, \dots, x_{n_i}) = \sum_{j=1}^{n_i} w_j \cdot \mathbf{E} = \left(\sum_{j=1}^{n_i} w_j \right) \cdot \mathbf{E}$$

- ▶ the sum over the word vectors in the document.

Word Embeddings can be used for any NLP Task

- ▶ For example, to predict political party from political speech:
 - ▶ could represent each document as concatenated embedding vectors for each word in the document.

Word Embeddings can be used for any NLP Task

- ▶ For example, to predict political party from political speech:
 - ▶ could represent each document as concatenated embedding vectors for each word in the document.
 - ▶ DNN will learn the optimal vector for each word in predicting the political party.

Practical Tip: Embedded Hashed N-Grams

Goldberg (2017) proposes the hashing vectorizer (`keras.preprocessing.text.hashing_trick`) as an efficient alternative to CNN's:

- ▶ Allocate $n_w \approx 1$ million rows to an embedding matrix \mathbf{E}
- ▶ Assign n-grams to embedding indexes with hashing function
- ▶ train MLP on top of embedding layer.

Practical Tip: Embedded Hashed N-Grams

Goldberg (2017) proposes the hashing vectorizer (`keras.preprocessing.text.hashing_trick`) as an efficient alternative to CNN's:

- ▶ Allocate $n_w \approx 1$ million rows to an embedding matrix \mathbf{E}
- ▶ Assign n-grams to embedding indexes with hashing function
- ▶ train MLP on top of embedding layer.
- ▶ Captures the local predictive power of n-grams without building vocabulary or costly training of CNN.

Outline

Intro to Neural Nets

Practicalities

Application to Fact-Value Distinction

Autoencoders

Embedding Layers

Word Embeddings

Word2Vec and GloVe

Word2Vec & GloVe

- ▶ “Word embeddings” often refers to Word2Vec or GloVe – these are particular (popular) models for producing word embeddings.
 - ▶ the goal: represent the meaning of words by the neighboring words – their **contexts**.

Word2Vec & GloVe

- ▶ “Word embeddings” often refers to Word2Vec or GloVe – these are particular (popular) models for producing word embeddings.
 - ▶ the goal: represent the meaning of words by the neighboring words – their **contexts**.
 - ▶ rather than predicting some metadata (such as citations) they predict the co-occurrence of neighboring words.

Word2Vec & GloVe

- ▶ “Word embeddings” often refers to Word2Vec or GloVe – these are particular (popular) models for producing word embeddings.
 - ▶ the goal: represent the meaning of words by the neighboring words – their **contexts**.
 - ▶ rather than predicting some metadata (such as citations) they predict the co-occurrence of neighboring words.
- ▶ “You shall know a word by the company it keeps”:
 - ▶ “He filled the **wampimuk**, passed it around and we all drunk some.”
 - ▶ “We found a little, hairy **wampimuk** sleeping behind the tree.”

Words and Contexts

A long line of NLP research aims to capture the distributional properties of words using a **word-context matrix** M :

- ▶ each row w represents a **word** (e.g. “income”), each column c represents a linguistic **context** in which words can occur (e.g. “corporate ___ tax”).
 - ▶ A matrix entry $M_{[w,c]}$ quantifies the strength of association between a word and a context in a large corpus.

Words and Contexts

A long line of NLP research aims to capture the distributional properties of words using a **word-context matrix M** :

- ▶ each row w represents a **word** (e.g. “income”), each column c represents a linguistic **context** in which words can occur (e.g. “corporate ____ tax”).
 - ▶ A matrix entry $M_{[w,c]}$ quantifies the strength of association between a word and a context in a large corpus.
- ▶ each word (row) w has a vector $M_{[w,:]}$ giving a distribution over contexts.
 - ▶ normally, these vectors have a spatial interpretation → geometric distances between word vectors reflect semantic distances between words.

Words and Contexts

A long line of NLP research aims to capture the distributional properties of words using a **word-context matrix M** :

- ▶ each row w represents a **word** (e.g. “income”), each column c represents a linguistic **context** in which words can occur (e.g. “corporate ____ tax”).
 - ▶ A matrix entry $M_{[w,c]}$ quantifies the strength of association between a word and a context in a large corpus.
- ▶ each word (row) w has a vector $M_{[w,:]}$ giving a distribution over contexts.
 - ▶ normally, these vectors have a spatial interpretation → geometric distances between word vectors reflect semantic distances between words.
 - ▶ Different definitions of contexts and different measures of association → different types of word vectors.

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get ($w = \text{"the"}$, $c = \text{"tabby"}$), ($w = \text{"tabby"}$, $c = \text{"the"}$), ($w = \text{"tabby"}$, $c = \text{"cat"}$), and ($w = \text{"cat"}$, $c = \text{"tabby"}$)

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get ($w = \text{"the"}$, $c = \text{"tabby"}$), ($w = \text{"tabby"}$, $c = \text{"the"}$), ($w = \text{"tabby"}$, $c = \text{"cat"}$), and ($w = \text{"cat"}$, $c = \text{"tabby"}$)
 - ▶ note the symmetry between words and contexts in this case.

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get $(w = \text{"the"}, c = \text{"tabby"})$, $(w = \text{"tabby"}, c = \text{"the"})$, $(w = \text{"tabby"}, c = \text{"cat"})$, and $(w = \text{"cat"}, c = \text{"tabby"})$
 - ▶ note the symmetry between words and contexts in this case.
- ▶ Could extend this to words within a window of two:
 - ▶ add $(w = \text{"the"}, c = \text{"cat"})$, $(w = \text{"cat"}, c = \text{"the"})$

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get $(w = \text{"the"}, c = \text{"tabby"})$, $(w = \text{"tabby"}, c = \text{"the"})$, $(w = \text{"tabby"}, c = \text{"cat"})$, and $(w = \text{"cat"}, c = \text{"tabby"})$
 - ▶ note the symmetry between words and contexts in this case.
- ▶ Could extend this to words within a window of two:
 - ▶ add $(w = \text{"the"}, c = \text{"cat"})$, $(w = \text{"cat"}, c = \text{"the"})$
 - ▶ word2vec and glove generally use 5- or 10-word windows.

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get $(w = \text{"the"}, c = \text{"tabby"})$, $(w = \text{"tabby"}, c = \text{"the"})$, $(w = \text{"tabby"}, c = \text{"cat"})$, and $(w = \text{"cat"}, c = \text{"tabby"})$
 - ▶ note the symmetry between words and contexts in this case.
- ▶ Could extend this to words within a window of two:
 - ▶ add $(w = \text{"the"}, c = \text{"cat"})$, $(w = \text{"cat"}, c = \text{"the"})$
 - ▶ word2vec and glove generally use 5- or 10-word windows.
- ▶ The context could be the tuple of preceding and subsequent words:
 - ▶ $(w = \text{"tabby"}, c = (\text{"the"}, \text{"cat"}))$

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get $(w = \text{"the"}, c = \text{"tabby"})$, $(w = \text{"tabby"}, c = \text{"the"})$, $(w = \text{"tabby"}, c = \text{"cat"})$, and $(w = \text{"cat"}, c = \text{"tabby"})$
 - ▶ note the symmetry between words and contexts in this case.
- ▶ Could extend this to words within a window of two:
 - ▶ add $(w = \text{"the"}, c = \text{"cat"})$, $(w = \text{"cat"}, c = \text{"the"})$
 - ▶ word2vec and glove generally use 5- or 10-word windows.
- ▶ The context could be the tuple of preceding and subsequent words:
 - ▶ $(w = \text{"tabby"}, c = (\text{"the"}, \text{"cat"}))$
- ▶ Could include all words in the same sentence.

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get $(w = \text{"the"}, c = \text{"tabby"})$, $(w = \text{"tabby"}, c = \text{"the"})$, $(w = \text{"tabby"}, c = \text{"cat"})$, and $(w = \text{"cat"}, c = \text{"tabby"})$
 - ▶ note the symmetry between words and contexts in this case.
- ▶ Could extend this to words within a window of two:
 - ▶ add $(w = \text{"the"}, c = \text{"cat"})$, $(w = \text{"cat"}, c = \text{"the"})$
 - ▶ word2vec and glove generally use 5- or 10-word windows.
- ▶ The context could be the tuple of preceding and subsequent words:
 - ▶ $(w = \text{"tabby"}, c = (\text{"the"}, \text{"cat"}))$
- ▶ Could include all words in the same sentence.
 - ▶ or same paragraph
 - ▶ or nouns in the same sentence
 - ▶ or syntactically connected words (from the parse tree)

Defining the context

- ▶ The simplest definition of context is neighboring words:
 - ▶ for “the tabby cat”: we get $(w = \text{"the"}, c = \text{"tabby"})$, $(w = \text{"tabby"}, c = \text{"the"})$, $(w = \text{"tabby"}, c = \text{"cat"})$, and $(w = \text{"cat"}, c = \text{"tabby"})$
 - ▶ note the symmetry between words and contexts in this case.
- ▶ Could extend this to words within a window of two:
 - ▶ add $(w = \text{"the"}, c = \text{"cat"})$, $(w = \text{"cat"}, c = \text{"the"})$
 - ▶ word2vec and glove generally use 5- or 10-word windows.
- ▶ The context could be the tuple of preceding and subsequent words:
 - ▶ $(w = \text{"tabby"}, c = (\text{"the"}, \text{"cat"}))$
- ▶ Could include all words in the same sentence.
 - ▶ or same paragraph
 - ▶ or nouns in the same sentence
 - ▶ or syntactically connected words (from the parse tree)
 - ▶ ...
- ▶ Etc.

Association Measures

- ▶ Let $\mathbf{M}_{[w,c]} = f_M(w, c)$ where w and c are lookups to words in the w vocabulary and c vocabulary.
 - ▶ “word” could also mean phrases or more complicated objects.

Association Measures

- ▶ Let $\mathbf{M}_{[w,c]} = f_M(w, c)$ where w and c are lookups to words in the w vocabulary and c vocabulary.
 - ▶ “word” could also mean phrases or more complicated objects.
- ▶ **counts:** $f_M(w, c) = \#(w, c)$, the number of times w appeared along with context c .
- ▶ **document frequencies:** $f_M(w, c) = \frac{\#(w, c)}{n_D}$

Association Measures

- ▶ Let $\mathbf{M}_{[w,c]} = f_M(w, c)$ where w and c are lookups to words in the w vocabulary and c vocabulary.
 - ▶ “word” could also mean phrases or more complicated objects.
- ▶ **counts:** $f_M(w, c) = \#(w, c)$, the number of times w appeared along with context c .
- ▶ **document frequencies:** $f_M(w, c) = \frac{\#(w, c)}{n_D}$
 - ▶ puts high weight on very common contexts that are shared across many words (e.g., “the cat” will be weighted higher than “tabby cat”)

Association Measures

- ▶ Let $\mathbf{M}_{[w,c]} = f_M(w, c)$ where w and c are lookups to words in the w vocabulary and c vocabulary.
 - ▶ “word” could also mean phrases or more complicated objects.
- ▶ **counts:** $f_M(w, c) = \#(w, c)$, the number of times w appeared along with context c .
- ▶ **document frequencies:** $f_M(w, c) = \frac{\#(w, c)}{n_D}$
 - ▶ puts high weight on very common contexts that are shared across many words (e.g., “the cat” will be weighted higher than “tabby cat”)
- ▶ **Point-wise mutual information (PMI)** normalized for high frequency:

$$(w, c) = \log \frac{\Pr(w, c)}{\Pr(w) \Pr(c)} = \log \frac{\frac{\#(w, c)}{n_D}}{\frac{\#(w)}{n_D} \frac{\#(c)}{n_D}} = \log \frac{n_D \#(w, c)}{\#(w) \#(c)}$$

where $\#(w)$ and $\#(c)$ are the corpus counts for w and c , respectively.

Issues with PMI

$$f_M(w, c) = \text{PMI}(w, c) = \log \frac{\#(w, c)}{\#(w)\#(c)}$$

- ▶ Issue 1: with $\text{PMI}(\cdot)$, unseen (w, c) pairs will take value $-\infty$.
 - ▶ So positive PMI (PPMI) is preferred:

$$\text{PPMI}(w, c) = \max(\text{PMI}(w, c), 0)$$

Issues with PMI

$$f_M(w, c) = \text{PMI}(w, c) = \log \frac{\#(w, c)}{\#(w)\#(c)}$$

- ▶ Issue 1: with $\text{PMI}(\cdot)$, unseen (w, c) pairs will take value $-\infty$.
 - ▶ So positive PMI (PPMI) is preferred:

$$\text{PPMI}(w, c) = \max(\text{PMI}(w, c), 0)$$

- ▶ Issue 2: PMI assigns high value to rare word-context pairs.
 - ▶ so impose a minimum count threshold on (w, c) pairs; below the threshold, set to zero.

\mathbf{M} is too high-dimensional

- ▶ \mathbf{M} is $n_w \times n_c$
 - ▶ if c is drawn from from the vocabulary of a reasonably large corpus, the associated word vectors $\{v_1 = \mathbf{M}_{[w_1, :]}, v_2 = \mathbf{M}_{[w_2, :]}, \dots\}$ are too high-dimensional to be useful.

\mathbf{M} is too high-dimensional

- ▶ \mathbf{M} is $n_w \times n_c$
 - ▶ if c is drawn from the vocabulary of a reasonably large corpus, the associated word vectors $\{v_1 = \mathbf{M}_{[w_1,:]}, v_2 = \mathbf{M}_{[w_2,:]}, \dots\}$ are too high-dimensional to be useful.
- ▶ The standard approach is singular value decomposition (SVD):
 - ▶ factorize $\mathbf{M} \in \mathbb{R}^{n_w \times n_c}$ into a word matrix $\mathbf{W} \in \mathbb{R}^{n_w \times n_E}$ and context matrix $\mathbf{C} \in \mathbb{R}^{n_c \times n_E}$
 - ▶ such that $\tilde{\mathbf{M}} = \mathbf{W}\mathbf{C}'$ is the best rank- n_E approximation of \mathbf{M} .

\mathbf{M} is too high-dimensional

- ▶ \mathbf{M} is $n_w \times n_c$
 - ▶ if c is drawn from the vocabulary of a reasonably large corpus, the associated word vectors $\{v_1 = \mathbf{M}_{[w_1,:]}, v_2 = \mathbf{M}_{[w_2,:]}, \dots\}$ are too high-dimensional to be useful.
- ▶ The standard approach is singular value decomposition (SVD):
 - ▶ factorize $\mathbf{M} \in \mathbb{R}^{n_w \times n_c}$ into a word matrix $\mathbf{W} \in \mathbb{R}^{n_w \times n_E}$ and context matrix $\mathbf{C} \in \mathbb{R}^{n_c \times n_E}$
 - ▶ such that $\tilde{\mathbf{M}} = \mathbf{W}\mathbf{C}'$ is the best rank- n_E approximation of \mathbf{M} .
 - ▶ $\tilde{\mathbf{M}}$ can be seen as a “smoothed” version of \mathbf{M} ; “missing” values are filled in, etc.

\mathbf{M} is too high-dimensional


- ▶ \mathbf{M} is $n_w \times n_c$
 - ▶ if c is drawn from the vocabulary of a reasonably large corpus, the associated word vectors $\{v_1 = \mathbf{M}_{[w_1,:]}, v_2 = \mathbf{M}_{[w_2,:]}, \dots\}$ are too high-dimensional to be useful.
- ▶ The standard approach is singular value decomposition (SVD):
 - ▶ factorize $\mathbf{M} \in \mathbb{R}^{n_w \times n_c}$ into a word matrix $\mathbf{W} \in \mathbb{R}^{n_w \times n_E}$ and context matrix $\mathbf{C} \in \mathbb{R}^{n_c \times n_E}$
 - ▶ such that $\tilde{\mathbf{M}} = \mathbf{W}\mathbf{C}'$ is the best rank- n_E approximation of \mathbf{M} .
 - ▶ $\tilde{\mathbf{M}}$ can be seen as a “smoothed” version of \mathbf{M} ; “missing” values are filled in, etc.
- ▶ \mathbf{W} is the matrix of word vectors (word embeddings):
 - ▶ relatively low-dimensional ($n_E \ll n_w$, typically between 50 and 300)
 - ▶ dense, rather than sparse.

M is too high-dimensional

- ▶ M is $n_w \times n_c$
 - ▶ if c is drawn from the vocabulary of a reasonably large corpus, the associated word vectors $\{v_1 = M_{[w_1, :]}, v_2 = M_{[w_2, :]}, \dots\}$ are too high-dimensional to be useful.
- ▶ The standard approach is singular value decomposition (SVD):
 - ▶ factorize $M \in \mathbb{R}^{n_w \times n_c}$ into a word matrix $W \in \mathbb{R}^{n_w \times n_E}$ and context matrix $C \in \mathbb{R}^{n_c \times n_E}$
 - ▶ such that $\tilde{M} = WC'$ is the best rank- n_E approximation of M .
 - ▶ \tilde{M} can be seen as a “smoothed” version of M ; “missing” values are filled in, etc.
- ▶ W is the matrix of word vectors (word embeddings):
 - ▶ relatively low-dimensional ($n_E \ll n_w$, typically between 50 and 300)
 - ▶ dense, rather than sparse.
 - ▶ similarity measures between rows of W approximate similarity measures between rows of M

Word Similarity

- ▶ Once words are represented as vectors $\{v_1 = \mathbf{M}_{[w_1,:]}, v_2 = \mathbf{M}_{[w_2,:]}, \dots\}$, we can use linear algebra to understand the relationships between words:
 - ▶ Words that are geometrically close to each other are similar: e.g. “dog” and “cat”:




- ▶ The standard metric for comparing vectors is cosine similarity:

$$\cos \theta = \frac{v_1 \cdot v_2}{||v_1|| ||v_2||}$$

- ▶ alternatives include e.g. Jaccard similarity (Goldberg 2017)

Word Similarity

- ▶ Once words are represented as vectors $\{v_1 = \mathbf{M}_{[w_1,:]}, v_2 = \mathbf{M}_{[w_2,:]}, \dots\}$, we can use linear algebra to understand the relationships between words:
 - ▶ Words that are geometrically close to each other are similar: e.g. “dog” and “cat”:



- ▶ The standard metric for comparing vectors is cosine similarity:

$$\cos \theta = \frac{v_1 \cdot v_2}{||v_1|| ||v_2||}$$

- ▶ alternatives include e.g. Jaccard similarity (Goldberg 2017)
- ▶ Thanks to linearity, can compute similarities between groups of words by averaging the groups.

Word2Vec

- ▶ When people mention “word2vec”, they are usually talking about a particular word-embedding model with good performance on a range of analogy and prediction tasks.

Word2Vec

- ▶ When people mention “word2vec”, they are usually talking about a particular word-embedding model with good performance on a range of analogy and prediction tasks.
- ▶ How does it learn the meaning of the word “fox”?
 - ▶ By comparing true instances of the word fox (“The quick brown **fox** jumps over the lazy dog”)
 - ▶ to fake (randomly sampled) ones (“The prescription of **fox** is advised for this diagnosis”)

Word2Vec

- ▶ When people mention “word2vec”, they are usually talking about a particular word-embedding model with good performance on a range of analogy and prediction tasks.
- ▶ How does it learn the meaning of the word “fox”?
 - ▶ By comparing true instances of the word fox (“The quick brown **fox** jumps over the lazy dog”)
 - ▶ to fake (randomly sampled) ones (“The prescription of **fox** is advised for this diagnosis”)
- ▶ Word2Vec learns embedding vectors for the target word (“fox”) and context words (neighbors of “fox”) to distinguish true from false samples.

Word2Vec Objective

- ▶ The dataset is a collection of context pairs indexed by i :
 - ▶ $y_i = 1$ means correct (it appeared in the corpus)
 - ▶ $y_i = 0$ means incorrect (it was randomly drawn).

Word2Vec Objective

- ▶ The dataset is a collection of context pairs indexed by i :
 - ▶ $y_i = 1$ means correct (it appeared in the corpus)
 - ▶ $y_i = 0$ means incorrect (it was randomly drawn).
 - ▶ This random drawing of incorrect pairs is called “**negative sampling**”. It is a form of unsupervised learning.

Word2Vec Objective

- ▶ The dataset is a collection of context pairs indexed by i :
 - ▶ $y_i = 1$ means correct (it appeared in the corpus)
 - ▶ $y_i = 0$ means incorrect (it was randomly drawn).
 - ▶ This random drawing of incorrect pairs is called “**negative sampling**”. It is a form of unsupervised learning.
- ▶ Let $\hat{y}_i(w, c; \theta)$ be the predicted probability of a correct pair. Word2Vec minimizes the binary cross-entropy

$$L(\theta) = - \sum_{i=1}^{n_D} [y_i \log \hat{y}_i(w, c; \theta) + [1 - y_i] \log(1 - \hat{y}_i(w, c; \theta))]$$

The neural net

We have an embedding matrix for target words, \mathbf{E}_w , and for context words, \mathbf{E}_c .

- ▶ Let $v_w(w) = \mathbf{E}_w \cdot w$ and $v_c(c) = \mathbf{E}_c \cdot c$ be the associated lookup functions.

The neural net

We have an embedding matrix for target words, \mathbf{E}_w , and for context words, \mathbf{E}_c .

- ▶ Let $v_w(w) = \mathbf{E}_w \cdot w$ and $v_c(c) = \mathbf{E}_c \cdot c$ be the associated lookup functions.

Embedding layer for instance i :

- ▶ $\mathbf{w}_i = v_w(w_i)$

The neural net

We have an embedding matrix for target words, \mathbf{E}_w , and for context words, \mathbf{E}_c .

- ▶ Let $v_w(w) = \mathbf{E}_w \cdot w$ and $v_c(c) = \mathbf{E}_c \cdot c$ be the associated lookup functions.

Embedding layer for instance i :

- ▶ $\mathbf{w}_i = v_w(w_i)$
- ▶ **Skip-gram** step:
 - ▶ draw one word c_i from i 's context window (n_{cw} words before or after w_i)
 - ▶ look up $\mathbf{c}_i = v_c(c_i)$

The neural net

We have an embedding matrix for target words, \mathbf{E}_w , and for context words, \mathbf{E}_c .

- ▶ Let $v_w(w) = \mathbf{E}_w \cdot w$ and $v_c(c) = \mathbf{E}_c \cdot c$ be the associated lookup functions.

Embedding layer for instance i :

- ▶ $\mathbf{w}_i = v_w(w_i)$
- ▶ **Skip-gram** step:
 - ▶ draw one word c_i from i 's context window (n_{cw} words before or after w_i)
 - ▶ look up $\mathbf{c}_i = v_c(c_i)$

The embedding $[\mathbf{w}; \mathbf{c}]$ is input to MLP1 (no activation) with sigmoid output:

$$\hat{y}(w, c) = \text{sigmoid}(([\mathbf{w}; \mathbf{c}] \cdot \omega_0) \cdot \omega_1)$$

The neural net

We have an embedding matrix for target words, \mathbf{E}_w , and for context words, \mathbf{E}_c .


- ▶ Let $v_w(w) = \mathbf{E}_w \cdot w$ and $v_c(c) = \mathbf{E}_c \cdot c$ be the associated lookup functions.

Embedding layer for instance i :

- ▶ $\mathbf{w}_i = v_w(w_i)$
- ▶ **Skip-gram** step:
 - ▶ draw one word c_i from i 's context window (n_{cw} words before or after w_i)
 - ▶ look up $\mathbf{c}_i = v_c(c_i)$

The embedding $[\mathbf{w}; \mathbf{c}]$ is input to MLP1 (no activation) with sigmoid output:

$$\hat{y}(w, c) = \text{sigmoid}(([\mathbf{w}; \mathbf{c}] \cdot \omega_0) \cdot \omega_1)$$



Does this relate to \mathbf{M} ?

- ▶ Word2Vec produces embedding matrices \mathbf{E}_w and \mathbf{E}_c .
 - ▶ generally, context embeddings are discarded after training.

Does this relate to \mathbf{M} ?

- ▶ Word2Vec produces embedding matrices \mathbf{E}_w and \mathbf{E}_c .
 - ▶ generally, context embeddings are discarded after training.
- ▶ If we take $\tilde{\mathbf{M}} = \mathbf{E}_w \mathbf{E}_c'$, Levy and Goldberg (2014) show that word2vec is equivalent to factorizing a matrix \mathbf{M} with items

$$\mathbf{M}_{[w,c]} = \text{PMI}(w, c) - \log a$$

where a is a constant calibrating the amount of negative sampling.

GloVe Embeddings

- ▶ Pennington et al (2014) (GloVe = Global Vectors) take a different approach:
 - ▶ that does not require a neural net
- ▶ Let $\#(w, c)$ be the co-occurrence counts within some window, e.g. 10 words.

GloVe Embeddings

- ▶ Pennington et al (2014) (GloVe = Global Vectors) take a different approach:
 - ▶ that does not require a neural net
- ▶ Let $\#(w, c)$ be the co-occurrence counts within some window, e.g. 10 words.
- ▶ Learn \mathbf{W} and \mathbf{C} (containing vectors \mathbf{w} and \mathbf{c} corresponding to w and c , respectively), to solve

$$\min_{\mathbf{W}, \mathbf{C}} \sum_{(w, c)} \phi(\#(w, c)) (\mathbf{w} \cdot \mathbf{c} - \log(\#(w, c)))^2$$

- ▶ $\phi(\cdot)$ is weighting function to down-weight frequent words.

GloVe Embeddings

- ▶ Pennington et al (2014) (GloVe = Global Vectors) take a different approach:
 - ▶ that does not require a neural net
- ▶ Let $\#(w, c)$ be the co-occurrence counts within some window, e.g. 10 words.
- ▶ Learn \mathbf{W} and \mathbf{C} (containing vectors \mathbf{w} and \mathbf{c} corresponding to w and c , respectively), to solve

$$\min_{\mathbf{W}, \mathbf{C}} \sum_{(w, c)} \phi(\#(w, c)) (\mathbf{w} \cdot \mathbf{c} - \log(\#(w, c)))^2$$

- ▶ $\phi(\cdot)$ is weighting function to down-weight frequent words.
- ▶ Minimizes **squared difference** between:
 - ▶ **dot product of word vectors**, $\mathbf{w} \cdot \mathbf{c}$
 - ▶ **empirical co-occurrence**, $\log(\#(w, c))$
- ▶ Intuitively: words that co-occur have high correlation (dot product)

Word Embeddings Encode Linguistic Relations

Word Embeddings Encode Linguistic Relations

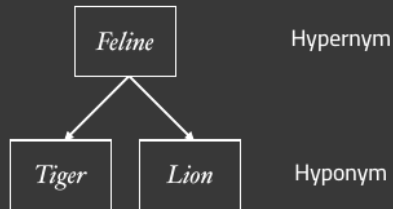
Synonymy



Antonymy



Hyponymy



Similarity vs. Relatedness (Budansky and Hirst, 2006)

- ▶ Semantic **similarity**: words sharing salient attributes / features
 - ▶ synonymy (car / automobile)
 - ▶ hypernymy (car / vehicle)
 - ▶ co-hyponymy (car / van / truck)

Similarity vs. Relatedness (Budansky and Hirst, 2006)

- ▶ Semantic **similarity**: words sharing salient attributes / features
 - ▶ synonymy (car / automobile)
 - ▶ hypernymy (car / vehicle)
 - ▶ co-hyponymy (car / van / truck)
- ▶ Semantic **relatedness**: words semantically associated without necessarily being similar
 - ▶ function (car / drive)
 - ▶ meronymy (car / tire)
 - ▶ location (car / road)
 - ▶ attribute (car / fast)

Similarity vs. Relatedness (Budansky and Hirst, 2006)

- ▶ Semantic **similarity**: words sharing salient attributes / features
 - ▶ synonymy (car / automobile)
 - ▶ hypernymy (car / vehicle)
 - ▶ co-hyponymy (car / van / truck)
- ▶ Semantic **relatedness**: words semantically associated without necessarily being similar
 - ▶ function (car / drive)
 - ▶ meronymy (car / tire)
 - ▶ location (car / road)
 - ▶ attribute (car / fast)
- ▶ Word embeddings will recover one or both of these relations, depending on how contexts and associated are constructed.

Most similar words to dog, depending on context window size

	2-word window	30-word window	
More paradigmatic		<u>kennel</u>	More syntagmatic
	cat	puppy	
	horse	pet	
	fox	bitch	
	pet	terrier	
	rabbit	rottweiler	
	pig	canine	
	animal	cat	
	mongrel	<u>bark</u>	
	sheep	alsatian	
	pigeon		

- ▶ Small windows pick up substitutable words; large windows pick up topics.

Target Word	BoW5	BoW2	Deps
batman	nightwing aquaman catwoman superman manhunter	superman superboy aquaman catwoman batgirl	superman superboy supergirl catwoman aquaman
hogwarts	dumbledore hallows half-blood malfoy snape	evernight sunnydale garderobe blandings collinwood	sunnydale collinwood calarts greendale millfield
turing	nondeterministic non-deterministic computability deterministic finite-state	non-deterministic finite-state nondeterministic buchi primality	Pauling hotelling hetting lessing hamming
florida	gainesville fla jacksonville tampa lauderdale	fla alabama gainesville tallahassee texas	texas louisiana georgia california carolina
object-oriented	aspect-oriented smalltalk event-driven prolog domain-specific	aspect-oriented event-driven objective-c dataflow 4gl	event-driven domain-specific rule-based data-driven human-centered
dancing	singing dance dances dancers tap-dancing	singing dance dances breakdancing clowning	singing rapping breakdancing miming busking

Parts of Speech and Phrases

- ▶ In the default model multiple senses of a word are merged.
 - ▶ e.g. “I like a bird” (verb) and “I am like a bird” (preposition).

Parts of Speech and Phrases

- ▶ In the default model multiple senses of a word are merged.
 - ▶ e.g. “I like a bird” (verb) and “I am like a bird” (preposition).
- ▶ Can significantly improve the quality of embeddings in these cases by attaching the POS to the word (e.g. “like:verb”, “like:prep”) before training.

Parts of Speech and Phrases

- ▶ In the default model multiple senses of a word are merged.
 - ▶ e.g. “I like a bird” (verb) and “I am like a bird” (preposition).
- ▶ Can significantly improve the quality of embeddings in these cases by attaching the POS to the word (e.g. “like:verb”, “like:prep”) before training.
- ▶ The default model only works by word, but “new york \neq ”new” + “york”
 - ▶ it makes sense to tokenize phrases together (see Week 2 lecture) before training.

The black sheep problem

- ▶ The trivial or obvious features of a word are not mentioned in standard corpora.

The black sheep problem

- ▶ The trivial or obvious features of a word are not mentioned in standard corpora.
- ▶ For example, although most sheep are white, you rarely see the phrase “white sheep”.
 - ▶ so word2vec tells you $\text{sim}(\text{black}, \text{sheep}) > \text{sim}(\text{white}, \text{sheep})$.

The black sheep problem

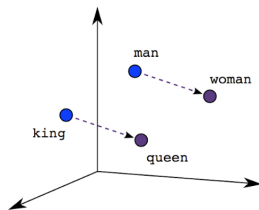
- ▶ The trivial or obvious features of a word are not mentioned in standard corpora.
- ▶ For example, although most sheep are white, you rarely see the phrase “white sheep”.
 - ▶ so word2vec tells you $\text{sim}(\text{black}, \text{sheep}) > \text{sim}(\text{white}, \text{sheep})$.
- ▶ This is really important when we will use embeddings to analysis social biases. And I don't see a solution to it.

The black sheep problem

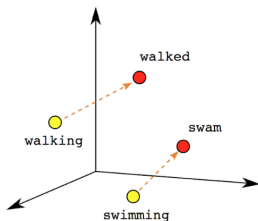
- ▶ The trivial or obvious features of a word are not mentioned in standard corpora.
- ▶ For example, although most sheep are white, you rarely see the phrase “white sheep”.
 - ▶ so word2vec tells you $\text{sim}(\text{black}, \text{sheep}) > \text{sim}(\text{white}, \text{sheep})$.
- ▶ This is really important when we will use embeddings to analysis social biases. And I don't see a solution to it.
- ▶ Relatedly, antonyms are often rated similarly, be careful with that.

Vector Directions \leftrightarrow Meaning


- ▶ Intriguingly, word2vec algebra can depict conceptual, analogical relationships between words:



Male-Female



Verb Tense



Country-Capital

Word Embeddings for Analogies

$$\text{vec}(\textit{king}) - \text{vec}(\textit{man}) + \text{vec}(\textit{woman}) \approx \text{vec}(\textit{queen})$$

Word Embeddings for Analogies

$$\text{vec}(\textit{king}) - \text{vec}(\textit{man}) + \text{vec}(\textit{woman}) \approx \text{vec}(\textit{queen})$$

- More generally: The analogy $a_1 : b_1 :: a_2 : b_2$ can be solved (that is, find b_2 given a_1, b_1, a_2) by

$$\arg \max_{b_2 \in V} \cos(b_2, a_2 - a_1 + b_1)$$

where V excludes (a_1, b_1, a_2) .

Word Embeddings for Analogies

$$\text{vec}(\textit{king}) - \text{vec}(\textit{man}) + \text{vec}(\textit{woman}) \approx \text{vec}(\textit{queen})$$

- More generally: The analogy $a_1 : b_1 :: a_2 : b_2$ can be solved (that is, find b_2 given a_1, b_1, a_2) by

$$\arg \max_{b_2 \in V} \cos(b_2, a_2 - a_1 + b_1)$$

where V excludes (a_1, b_1, a_2) .

- Often works better with normalized vectors (so that one long vector doesn't wash out the others)

Word Embeddings for Analogies

$$\text{vec}(\text{king}) - \text{vec}(\text{man}) + \text{vec}(\text{woman}) \approx \text{vec}(\text{queen})$$

- ▶ More generally: The analogy $a_1 : b_1 :: a_2 : b_2$ can be solved (that is, find b_2 given a_1, b_1, a_2) by

$$\arg \max_{b_2 \in V} \cos(b_2, a_2 - a_1 + b_1)$$

where V excludes (a_1, b_1, a_2) .

- ▶ Often works better with normalized vectors (so that one long vector doesn't wash out the others)
- ▶ Levy and Goldberg (2014) recommend the following “CosMul” metric which tends to perform better:

$$\arg \max_{b_2 \in V} \frac{\cos(b_2, a_2) \cos(b_2, b_1)}{\cos(b_2, a_1) + \epsilon}$$

- ▶ requires normalized, non-negative vectors (can transform using $(x+1)/2$)
- ▶ ϵ is a small smoothing parameter.

Word embeddings vs topic models

Ben Schmidt:

- ▶ Topic models reduce words to core meanings to understand documents more clearly.

Word embeddings vs topic models

Ben Schmidt:

- ▶ Topic models reduce words to core meanings to understand documents more clearly.
- ▶ Word embedding models ignore information about individual documents to better understand the relationships between words.

Tokenizing for Embeddings

- ▶ embeddings work better with more information about the placement of words in sentences.
 - ▶ don't drop stopwords/function-words
 - ▶ should include tokens for start of sentence and end of sentence
 - ▶ should include a special token for out-of-vocabulary words
 - ▶ or replace with the part of speech tag

Word Dropout

- ▶ When training models, words can be randomly replaced with the unknown symbol with some small probability (Iyer et al 2015).
- ▶ Prevents models from relying too much on particular words.

K-means clustering with Word Embeddings

Income Tax (Pensions Topic and Health Care Topic)

pension_board have_attain_the_age
in_excess_of_year
retir_purpos
such_depend
such_servic
biweek_pay_period

medicar servic in accord vocat rehabilit
legal_settlement
admiss_center
self-support
depend_children
coron medic condit
babiday_servic
cerebr_palsi

Sales Tax (Retail Topic and Health Care Topic)

fuel_dealer
retail_store
such_distributor
har_race

psychiatrist juvenil_offend
state_plan
educ_or_train
retard_servic
aid_to_famili
cost_of_health
first_aid

- ▶ Clustered phrases affecting tax revenues (Ash 2018); Green words tend to increase revenues; red words tend to decrease revenues.

Word Mover Distance

- ▶ Cosine distance treats synonyms as just as close as totally unrelated words.

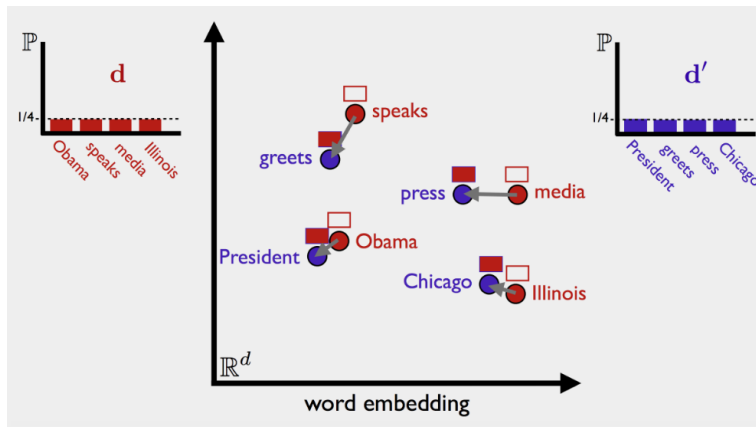
Word Mover Distance

- ▶ Cosine distance treats synonyms as just as close as totally unrelated words.
- ▶ Word mover distance between two texts is given by:
 - ▶ total amount of “mass” needed to move words from one side into the other
 - ▶ multiplied by the distance the words need to move
 - ▶ Kusner, Sun, Kolkin, and Weinberger (ICML 2015)

Word Mover Distance

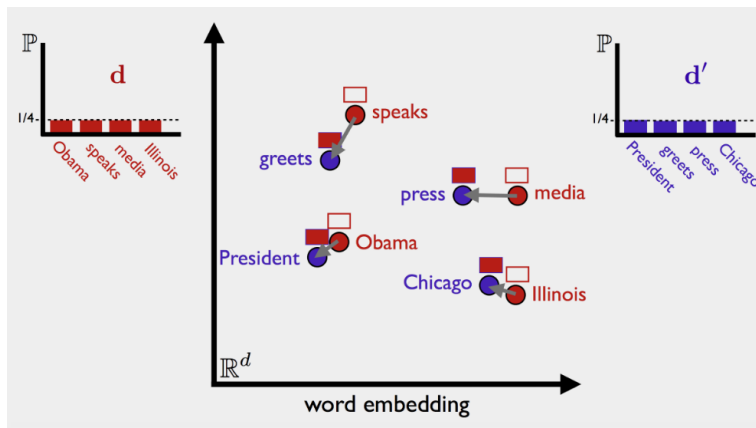
- ▶ Cosine distance treats synonyms as just as close as totally unrelated words.
- ▶ Word mover distance between two texts is given by:
 - ▶ total amount of “mass” needed to move words from one side into the other
 - ▶ multiplied by the distance the words need to move
 - ▶ Kusner, Sun, Kolkin, and Weinberger (ICML 2015)
- ▶ Requires measure of distance between words (word embeddings).
 - ▶ see wmd package in Python.

Illustration



- ▶ d (obama speaks media illinois) is orthogonal to d' (president greet press chicago):
 - ▶ cosine similarity is zero

Illustration



- ▶ d (obama speaks media illinois) is orthogonal to d' (president greet press chicago):
 - ▶ cosine similarity is zero
 - ▶ Word mover distance sums the shortest distances between the words in the documents.

Pre-trained word embeddings

- ▶ For some NLP tasks, you might not need to train your own vectors.
 - ▶ or your corpus might be too small to do so.

Pre-trained word embeddings

- ▶ For some NLP tasks, you might not need to train your own vectors.
 - ▶ or your corpus might be too small to do so.
- ▶ In these cases, you can use a pre-trained model, like spaCy's GloVe embeddings:
 - ▶ one million vocabulary entries
 - ▶ 300-dimensional vectors
 - ▶ trained on the Common Crawl corpus

Pre-trained word embeddings

- ▶ For some NLP tasks, you might not need to train your own vectors.
 - ▶ or your corpus might be too small to do so.
- ▶ In these cases, you can use a pre-trained model, like spaCy's GloVe embeddings:
 - ▶ one million vocabulary entries
 - ▶ 300-dimensional vectors
 - ▶ trained on the Common Crawl corpus
- ▶ Can initialize prediction model using pre-trained embeddings.

Tips for using pre-trained embeddings

- ▶ Split training in two steps:
 - ▶ in first run, train the model with the first layer (the pre-trained embeddings) frozen.

Tips for using pre-trained embeddings

- ▶ Split training in two steps:
 - ▶ in first run, train the model with the first layer (the pre-trained embeddings) frozen.
 - ▶ in second run, un-freeze the embedding layer for fine tuning.