

# Sequencing Legal DNA

## NLP for Law and Political Economy

### 5. Neural Nets and Word Embeddings

# Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce  $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$ , a vector of predicted probabilities across classes for each document  $i$ .

# Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce  $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$ , a vector of predicted probabilities across classes for each document  $i$ .
- ▶ This vector is a compressed representation of the outcome-predictive text features  $\mathbf{x}_i$

# Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce  $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$ , a vector of predicted probabilities across classes for each document  $i$ .
- ▶ This vector is a compressed representation of the outcome-predictive text features  $\mathbf{x}_i$ 
  - ▶  $\mathbf{x}_i$  is itself a compressed representation of the unprocessed document  $\mathcal{D}_i$ .

# Machine Learning Produces Representations of the Data

- ▶ Text classifiers produce  $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \hat{\theta})$ , a vector of predicted probabilities across classes for each document  $i$ .
- ▶ This vector is a compressed representation of the outcome-predictive text features  $\mathbf{x}_i$ 
  - ▶  $\mathbf{x}_i$  is itself a compressed representation of the unprocessed document  $\mathcal{D}_i$ .
- ▶ Correspondingly: the parameters  $\hat{\theta}$  can also be understood as a compressed (or “learned”) representation:
  - ▶ it contains information about the training corpus, the text features, and the outcomes.

## Information in $\hat{\theta}$

- ▶ Say we train a multinomial logistic regression on a bag-of-words representation of the documents.
- ▶ Let  $\theta$  be the learned matrix of parameters relating words to outcomes:
  - ▶ It contains  $n_y$  columns =  $n_x$ -vectors representing the outcome classes.

## Information in $\hat{\theta}$

- ▶ Say we train a multinomial logistic regression on a bag-of-words representation of the documents.
- ▶ Let  $\theta$  be the learned matrix of parameters relating words to outcomes:
  - ▶ It contains  $n_y$  columns =  $n_x$ -vectors representing the outcome classes.
  - ▶ It contains  $n_x$  rows =  $n_y$ -vectors representing each word in the vocabulary.

## Information in $\hat{\theta}$

- ▶ Say we train a multinomial logistic regression on a bag-of-words representation of the documents.
- ▶ Let  $\theta$  be the learned matrix of parameters relating words to outcomes:
  - ▶ It contains  $n_y$  columns =  $n_x$ -vectors representing the outcome classes.
  - ▶ It contains  $n_x$  rows =  $n_y$ -vectors representing each word in the vocabulary.
- ▶ How to use this?
  - ▶ could cluster column vectors to understand which outcomes are similar/related.
  - ▶ could cluster row vectors to understand which features are similar/related.



## Preview of Word Embeddings

- ▶ Let's say  $\mathbf{x}_i$  is a bag-of-words representation for document  $i$  with length  $n_i$ . We can write

$$\mathbf{x}_i = \frac{1}{n_i} \sum_{l=1}^{n_i} \mathbf{x}_i^{[l]}$$

- ▶  $l$  indexes words in the the document
- ▶ each vector  $\mathbf{x}_i^{[l]}$  is an  $n_x$ -dimensional one-hot vector – all entries are zero except the single entry corresponding to the word at  $l$ , which is 1.

## Preview of Word Embeddings

- ▶ Let's say  $\mathbf{x}_i$  is a bag-of-words representation for document  $i$  with length  $n_i$ . We can write

$$\mathbf{x}_i = \frac{1}{n_i} \sum_{l=1}^{n_i} \mathbf{x}_i^{[l]}$$

- ▶  $l$  indexes words in the the document
- ▶ each vector  $\mathbf{x}_i^{[l]}$  is an  $n_x$ -dimensional one-hot vector – all entries are zero except the single entry corresponding to the word at  $l$ , which is 1.
- ▶ Now let  $\theta^{[l]}$  be the row of  $\theta$  corresponding to the word  $w_l$ . We can write

$$\hat{\mathbf{y}}_i = \frac{1}{n_i} \sum_{l=1}^{n_i} \theta^{[l]}$$

the sum of the  $n_y$ -dimensional word representations (the row vectors from above).

- ▶ this is called the “continuous bag of words (CBOW)” representation.
- ▶  $\theta$  is a word embedding matrix.

# Outline

Intro to Neural Nets

Practicalities

Autoencoders

Embedding Layers

Word Embeddings

Kozlowski, Evans, and Taddy (2019)

# “Neural Networks”

- ▶ “Neural”:
  - ▶ nothing like brains

# “Neural Networks”

- ▶ “Neural”:
  - ▶ nothing like brains
- ▶ “Networks”:
  - ▶ nothing to do with “networks” as normally understood – in particular, nothing to do with network theory in math or social science.

## Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.

## Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.

## Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.
- ▶ Training algorithms have improved – small tweaks have made a huge impact.



## Recent History

- ▶ NNs frequently outperform other ML techniques on large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.
- ▶ Training algorithms have improved – small tweaks have made a huge impact.
- ▶ Some theoretical limitations of NNs have turned out to be benign in practice – for example, they work well on non-convex functions.

# Will it last?

- ▶ Three key principles of deep learning will persist:
  - ▶ **Simplicity**
    - ▶ feature engineering is obsolete
    - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.

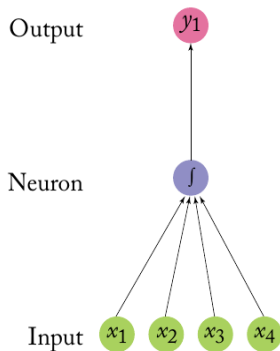
# Will it last?

- ▶ Three key principles of deep learning will persist:
  - ▶ **Simplicity**
    - ▶ feature engineering is obsolete
    - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.
  - ▶ **Scalability**
    - ▶ amenable to parallelization on GPUs or TPUs (tensor processing units)
    - ▶ trained on batches of data, so can be scaled to datasets of arbitrary size.

# Will it last?

- ▶ Three key principles of deep learning will persist:
  - ▶ **Simplicity**
    - ▶ feature engineering is obsolete
    - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.
  - ▶ **Scalability**
    - ▶ amenable to parallelization on GPUs or TPUs (tensor processing units)
    - ▶ trained on batches of data, so can be scaled to datasets of arbitrary size.
  - ▶ **Versatility and reusability**
    - ▶ can be trained on additional data without restarting from scratch, therefore amenable for continuous online learning.
    - ▶ deep-learning models are repurposable and thus reusable

## A “Neuron”



- ▶ A neuron multiplies each input by its weight, sums them, applies a non-linear function to the result, and passes the output.
  - ▶ e.g., the  $\int$  shape indicates a sigmoid transformation.

## In Notation

- ▶ The simplest neural network is called a perceptron:

$$\text{MLP0}(\mathbf{x}) = \mathbf{x} \cdot \boldsymbol{\omega}$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega} \in \mathbb{R}^{n_x \times n_y}$$

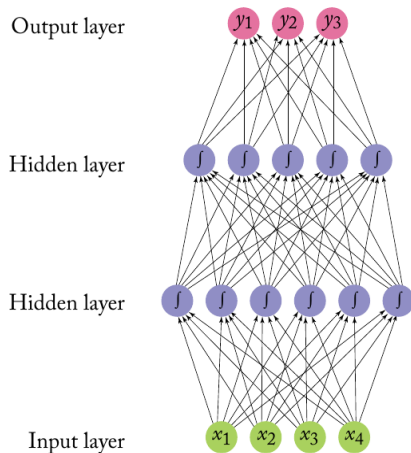
here,  $\boldsymbol{\omega}$  is the matrix of weights in the layer.

- ▶ In more standard notation, there would be an additional constant (or “bias”) term:

$$\text{MLP0}(\mathbf{x}) = \alpha + \mathbf{x} \cdot \boldsymbol{\omega}$$

- ▶ We leave it out by assuming that  $\mathbf{x}$  is de-meaned or has an extra column of ones.

# A Feed-Forward Neural Network



- A feed-forward network is simply a stack of linear models, separated by non-linear functions.

# Multi-Layer Perceptron

- ▶ An multi-layer perceptron (MLP) with one hidden layer is

$$\text{MLP1}(\mathbf{x}) = \mathbf{g}(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_y},$$

- ▶  $n_1$  = dimensionality in first (and only) hidden layer
- ▶  $\boldsymbol{\omega}_1$  = set of learnable weights for the first linear transformation of the inputs.
- ▶  $\mathbf{g}(\cdot)$  = an element-wise non-linear function (an “activation function”)
- ▶  $\boldsymbol{\omega}_2$  = weights on the second linear transformation leading to the output.



# Multi-Layer Perceptron

- ▶ An multi-layer perceptron (MLP) with one hidden layer is

$$\text{MLP1}(\mathbf{x}) = \mathbf{g}(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_y},$$

- ▶  $n_1$  = dimensionality in first (and only) hidden layer
- ▶  $\boldsymbol{\omega}_1$  = set of learnable weights for the first linear transformation of the inputs.
- ▶  $\mathbf{g}(\cdot)$  = an element-wise non-linear function (an “activation function”)
- ▶  $\boldsymbol{\omega}_2$  = weights on the second linear transformation leading to the output.
- ▶ MLP1 can approximate any continuous function on a closed and bounded subset of  $\mathbb{R}^n$ , and any mapping from one finite discrete space to another finite discrete space (Hornik et al 1989, Cybenko 1989).
  - ▶ But MLP1 would have to be exponentially large in some cases (Telgarsky 2016) .

## Two hidden layers

- ▶ Adding a second hidden layer gives

$$\text{MLP2}(\mathbf{x}) = \mathbf{g}_2(\mathbf{g}_1(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2) \cdot \boldsymbol{\omega}_3$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_2}, \boldsymbol{\omega}_3 \in \mathbb{R}^{n_2 \times n_y}$$

- ▶  $n_2$  = number of neurons in second hidden layer.

## Two hidden layers

- ▶ Adding a second hidden layer gives

$$\text{MLP2}(\mathbf{x}) = \mathbf{g}_2(\mathbf{g}_1(\mathbf{x} \cdot \boldsymbol{\omega}_1) \cdot \boldsymbol{\omega}_2) \cdot \boldsymbol{\omega}_3$$

$$\mathbf{x} \in \mathbb{R}^{n_x}, \boldsymbol{\omega}_1 \in \mathbb{R}^{n_x \times n_1}, \boldsymbol{\omega}_2 \in \mathbb{R}^{n_1 \times n_2}, \boldsymbol{\omega}_3 \in \mathbb{R}^{n_2 \times n_y}$$

- ▶  $n_2$  = number of neurons in second hidden layer.
- ▶ MLP2 can be written in the following decomposed notation:

$$\text{MLP2}(\mathbf{x}) =$$

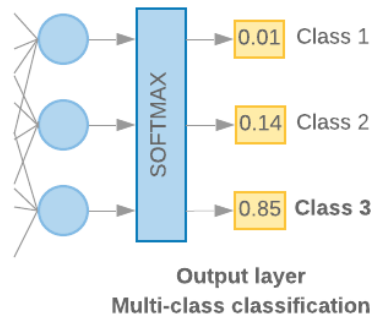
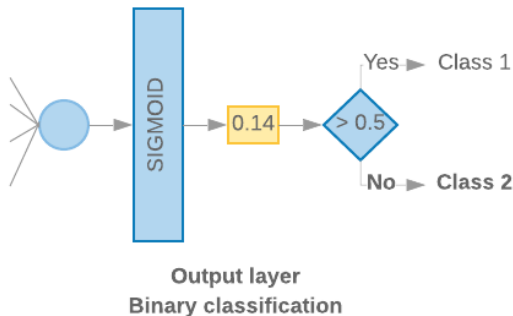
$$\mathbf{h}_1 = \mathbf{g}_1(\mathbf{x} \cdot \boldsymbol{\omega}_1)$$

$$\mathbf{h}_2 = \mathbf{g}_2(\mathbf{h}_1 \cdot \boldsymbol{\omega}_2)$$

$$\mathbf{y} = \mathbf{h}_2 \cdot \boldsymbol{\omega}_3$$

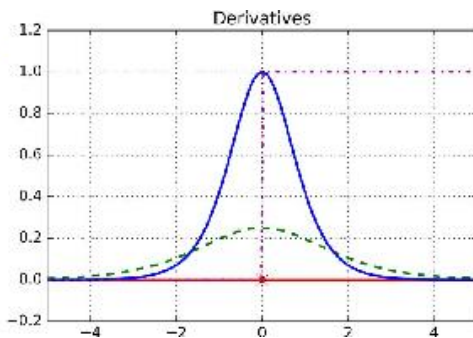
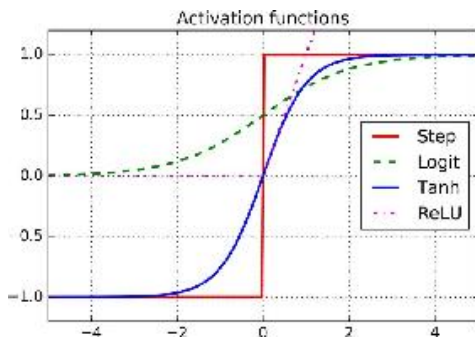
where  $\mathbf{h}_l$  give hidden layers.

## Constructing the Last Layer



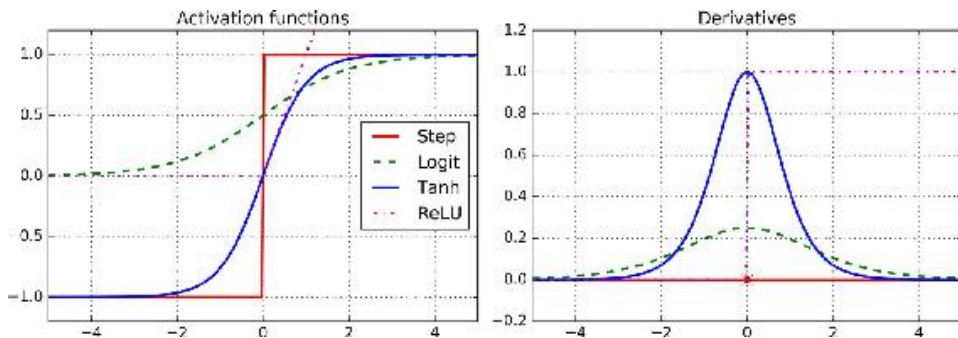
- ▶ MLPs will output a probability distribution across output classes.
  - ▶ can also output a real number, which would make a regression model.

## What to pick for $g(\cdot)$



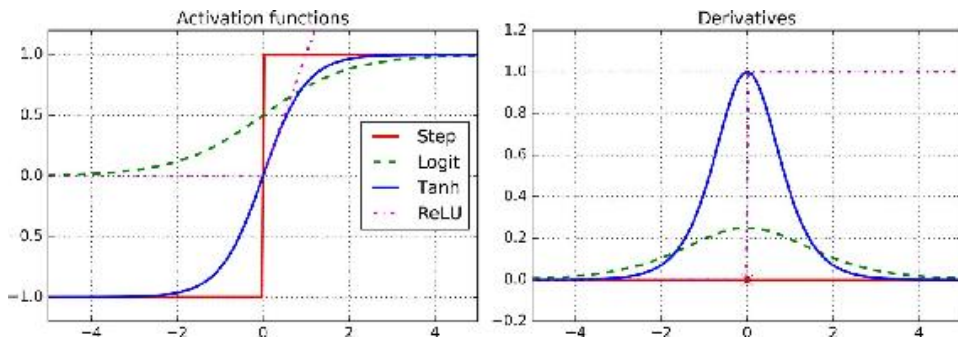
- logistic function:  $\text{logit}(z) = \frac{1}{1+\exp(-z)}$

## What to pick for $g(\cdot)$



- ▶ logistic function:  $\text{logit}(z) = \frac{1}{1+\exp(-z)}$
- ▶ hyperbolic tangent function:  $\tanh(z) = 2\sigma(2z) - 1$ 
  - ▶ ranges between -1 and 1 (rather than between 0 and 1, as the case with the logistic)
  - ▶ centered on zero, can speed up convergence

# What to pick for $g(\cdot)$



- ▶ logistic function:  $\text{logit}(z) = \frac{1}{1+\exp(-z)}$
- ▶ hyperbolic tangent function:  $\tanh(z) = 2\sigma(2z) - 1$ 
  - ▶ ranges between -1 and 1 (rather than between 0 and 1, as the case with the logistic)
  - ▶ centered on zero, can speed up convergence
- ▶ ReLU (rectified linear unit) function:  $\max\{0, z\}$ ,
  - ▶ deceptively simple, fast to compute, and very effective in practice
  - ▶ gradient does not saturate to zero for large values (but is flat below zero)

# Outline

Intro to Neural Nets

Practicalities

Autoencoders

Embedding Layers

Word Embeddings

Kozlowski, Evans, and Taddy (2019)



# MLP baseline for Text Classification

Google Developers Advice

1. Calculate the number of samples/number of words per sample ratio.
2. If this ratio is less than 1500, tokenize the text as n-grams and use a simple multi-layer perceptron (MLP) model to classify them.
  - ▶ In the case of N-grams models, Google testers found that MLPs tended to out-perform logistic regression and gradient boosting machines.

# Keras Basics

- ▶ See the Geron book and sample notebooks for Keras examples.
- ▶ “Dense” layer is the DNN baseline – means that all neurons are connected.

# Keras Basics

- ▶ See the Geron book and sample notebooks for Keras examples.
- ▶ “Dense” layer is the DNN baseline – means that all neurons are connected.
- ▶ Output layer:
  - ▶ for regression, do not use an activation function
  - ▶ for binary classification, use `activation='sigmoid'`
  - ▶ for multi-class classification, use `activation='softmax'`

# Loss function and metrics

- ▶ Loss function:
  - ▶ for regression, use `mean_squared_error`
  - ▶ for binary classification, use `binary_crossentropy`
  - ▶ for multi-class classification, use `sparse_categorical_crossentropy`

# Loss function and metrics

- ▶ Loss function:
  - ▶ for regression, use `mean_squared_error`
  - ▶ for binary classification, use `binary_crossentropy`
  - ▶ for multi-class classification, use `sparse_categorical_crossentropy`
- ▶ Metrics:
  - ▶ for classification, can use accuracy and  $F_1$
  - ▶ for regression, use  $R^2$

# Tuning NN Hyperparameters

- ▶ Number of hidden layers:
  - ▶ having a single hidden layer will generally give decent results.
    - ▶ more layers with fewer neurons can recover hierarchical relations and complex functions
    - ▶ for text classification, try one or two hidden layers as a baseline.

# Tuning NN Hyperparameters

- ▶ Number of hidden layers:
  - ▶ having a single hidden layer will generally give decent results.
    - ▶ more layers with fewer neurons can recover hierarchical relations and complex functions
    - ▶ for text classification, try one or two hidden layers as a baseline.
- ▶ Number of neurons:
  - ▶ a common practice is to set neuron counts like a funnel, with fewer and fewer neurons at each level
  - ▶ or just pick 128 neurons per layer
  - ▶ overall, better to have too many neurons, and use regularization

# Tuning NN Hyperparameters

- ▶ Number of hidden layers:
  - ▶ having a single hidden layer will generally give decent results.
    - ▶ more layers with fewer neurons can recover hierarchical relations and complex functions
    - ▶ for text classification, try one or two hidden layers as a baseline.
- ▶ Number of neurons:
  - ▶ a common practice is to set neuron counts like a funnel, with fewer and fewer neurons at each level
  - ▶ or just pick 128 neurons per layer
  - ▶ overall, better to have too many neurons, and use regularization
- ▶ Activation functions:
  - ▶ use ReLU in hidden layers



## Xavier and He Initialization

Activation function	Uniform distribution $[-r, r]$	Normal distribution
Logistic	$r = \sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
Hyperbolic tangent	$r = 4\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = 4\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$
ReLU (and its variants)	$r = \sqrt{2}\sqrt{\frac{6}{n_{\text{inputs}} + n_{\text{outputs}}}}$	$\sigma = \sqrt{2}\sqrt{\frac{2}{n_{\text{inputs}} + n_{\text{outputs}}}}$

- Connection weights should be initialized randomly according to a uniform distribution or normal distribution, as indicated in the table (see Geron Chapter 11).

## Other Activation Functions


- ▶ Leaky ReLU

$$\max(\alpha z, z)$$

where  $\alpha$  is set to a small number, such as .01, or learned in training.

- ▶ Exponential linear unit

$$\text{ELU}(z) = \begin{cases} \alpha(\exp(z) - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$



- ▶ Until recently, ELU has had the best performance so far, but it is slower than ReLU.

# Batch normalization

- ▶ Another trick to speed up training:
  - ▶ in between layers, zero-center and normalize the inputs to variance one.
  - ▶ normally done before a non-linear activation function

# Regularization for Sparse Models

- ▶ As with linear models, neural network parameters can be regularized with an L1 and/or L2 penalty to push weak neurons to zero and produce a sparse model.
- ▶ But usually its better/simpler to use dropout.

# Dropout

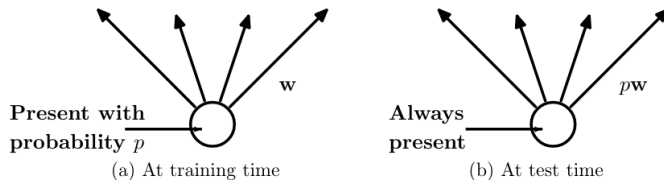


Figure 2: **Left:** A unit at training time that is present with probability  $p$  and is connected to units in the next layer with weights  $\mathbf{w}$ . **Right:** At test time, the unit is always present and the weights are multiplied by  $p$ . The output at test time is same as the expected output at training time.

Source: Srivastava et al, JMLR 2014

An elegant regularization technique:

- ▶ at every training step, every neuron has some probability (typically  $p = 0.5$ ) of being temporarily dropped out, so that it will be ignored at this step.
- ▶ at test time, neurons don't get dropped any more but coefficients are down-weighted by  $p$ .

# Dropout

- ▶ Approximately equivalent to averaging the output of  $N$  models (where  $N$  is the number of neurons).

# Dropout

- ▶ Approximately equivalent to averaging the output of  $N$  models (where  $N$  is the number of neurons).
- ▶ Neurons trained with dropout:
  - ▶ cannot co-adapt with neighboring neurons and must be independently useful.
  - ▶ cannot rely excessively on just a few input neurons; they have to pay attention to all input neurons.
  - ▶ makes the model less sensitive to slight changes in the inputs.

# Optimizers

- ▶ Choice of optimization algorithm is the topic of active research, which has shown that it can have a big impact on model performance.
  - ▶ Until recently, a good starting choice would be Adam (adaptive moment estimation), which is fast and usually works well. For robustness, can also try SGD.
  - ▶ A recent paper says that AdaBound dominates Adam or SGD.



## Early stopping

- ▶ A popular/efficient regularization method is to continually evaluate your model at regular intervals, and then to stop training when the test-set accuracy starts to decrease.

## Early stopping

- ▶ A popular/efficient regularization method is to continually evaluate your model at regular intervals, and then to stop training when the test-set accuracy starts to decrease.
- ▶ Split data into three sets: training, validation, and test.
  - ▶ every few epochs, check accuracy in validation set.
  - ▶ if it has gone down since last check, stop and use the model at the previous checkpoint.

## Batch Training with Large Data

- ▶ If data sets don't fit in memory, can load the data in batches from disk.
- ▶ can also continuously update a saved model.

# Outline

Intro to Neural Nets

Practicalities



**Autoencoders**

Embedding Layers

Word Embeddings


Kozlowski, Evans, and Taddy (2019)

## Remember the Swiss roll?




- The dimension reduction process matters: projecting down to two dimensions directly (left panel) might not isolate the variation we are interested in (as done in the right panel, which unrolls the Swiss Roll)


# Autoencoders: Domain-specific dimension reduction



- ▶ “Autoencoder” refers to a class of deep neural network that performs domain-specific dimension reduction.
  - ▶ They learn efficient encodings of the data, which can then be decoded back to a (minimally) lossy representation of the original data.
  - ▶ Can also randomly generate new data that looks like the training data.



(a) The architecture of SAE with  $(S-1)$  hidden layers in both encoder and decoder.



(b) The graph representation of SAE with  $(S-1)$  hidden layers in both encoder and decoder.

- ▶ Autoencoders work by stacking layers that gradually decrease in dimensionality to create the compressed representation ( $Z$ ), and then gradually increase in dimensionality to try to reconstruct the input.
  - ▶ the autoencoder is implicitly solving the problem of maximizing entropy in the bottleneck layer.

# Autoencoding for data visualization

- ▶ For 2D visualization, t-SNE is probably the best algorithm
  - ▶ but quite slow, and typically requires relatively low-dimensional data.



# Autoencoding for data visualization

- ▶ For 2D visualization, t-SNE is probably the best algorithm
  - ▶ but quite slow, and typically requires relatively low-dimensional data.
- ▶ Decent baseline:
  - ▶ use an autoencoder to compress your data to relatively low dimension (e.g. 32 dimensions)
  - ▶ then use t-SNE for mapping the compressed data to a 2D plane.

# Outline

Intro to Neural Nets

Practicalities

Autoencoders

Embedding Layers

Word Embeddings

Kozlowski, Evans, and Taddy (2019)

# What is an Embedding?

- ▶ A (relatively) low-dimensional vector representation of a categorical variable.
  - ▶ spatial location of the vector encodes predictive information about the category.

# What is an Embedding?

- ▶ A (relatively) low-dimensional vector representation of a categorical variable.
  - ▶ spatial location of the vector encodes predictive information about the category.
- ▶ e.g., trying to predict how employment responds to economic growth with data from U.S. states:
  - ▶ instead of including a fifty-dimensional categorical variable, include two-dimensional latitude and longitude

# What is an Embedding?

- ▶ A (relatively) low-dimensional vector representation of a categorical variable.
  - ▶ spatial location of the vector encodes predictive information about the category.
- ▶ e.g., trying to predict how employment responds to economic growth with data from U.S. states:
  - ▶ instead of including a fifty-dimensional categorical variable, include two-dimensional latitude and longitude
  - ▶ or initialize each state to a random two-dimensional vector, and let the model decide where to move the states to improve prediction on your task (e.g. ).

# An embedding layer is just matrix multiplication

- ▶ An embedding layer can be represented as

$$\underbrace{x}_{n_E \times 1} = \underbrace{\Omega}_{n_E \times n_w} \cdot \underbrace{w}_{n_w \times 1}$$

- ▶  $w$ , a categorical variable (e.g., representing a word)
  - ▶ one-hot vector with a single item equaling one.
  - ▶ The input to the embedding layer.

# An embedding layer is just matrix multiplication

- ▶ An embedding layer can be represented as

$$\underbrace{x}_{n_E \times 1} = \underbrace{\Omega}_{n_E \times n_w} \cdot \underbrace{w}_{n_w \times 1}$$

- ▶  $w$ , a categorical variable (e.g., representing a word)
  - ▶ one-hot vector with a single item equaling one.
  - ▶ The input to the embedding layer.
- ▶  $x$ , a dense representation of the variable.
  - ▶ The output of the embedding layer.

# An embedding layer is just matrix multiplication

- ▶ An embedding layer can be represented as


$$\underbrace{x}_{n_E \times 1} = \underbrace{\Omega}_{n_E \times n_w} \cdot \underbrace{w}_{n_w \times 1}$$


- ▶  $w$ , a categorical variable (e.g., representing a word)
  - ▶ one-hot vector with a single item equaling one.
  - ▶ The input to the embedding layer.
- ▶  $x$ , a dense representation of the variable.
  - ▶ The output of the embedding layer.
- ▶ An embedding matrix  $\Omega$ , learnable by the DNN




# The Embedding Matrix $\Omega$

- ▶ The model learns the weights of the embedding matrix in the same way that it would learn any model parameters.
- ▶ The rows of the matrix correspond to vectors for the  $n_w$  categories.
  - ▶ These are the “word vectors” that people talk about when they mention word embeddings or Word2Vec.







# Embedding Layers versus Dense Layers

- ▶ An embedding layer is statistically equivalent to a fully-connected dense layer with sparse data set as input and linear activation.

# Embedding Layers versus Dense Layers

- ▶ An embedding layer is statistically equivalent to a fully-connected dense layer with sparse data set as input and linear activation.
- ▶ Why use an embedding layer rather than a dense layer?
  - ▶ embedding layers are much faster for this purpose
  - ▶ batch updating with regularization and dropout do not work well on sparse data.

# Outline

Intro to Neural Nets

Practicalities

Autoencoders

Embedding Layers

**Word Embeddings**

Kozlowski, Evans, and Taddy (2019)

# Word Embeddings

- ▶ Embedding layer maps word indexes to dense vectors.



# Word Embeddings

- ▶ Embedding layer maps word indexes to dense vectors.
- ▶ Documents are lists of word indexes  $\{w_1, w_2, \dots, w_{n_i}\}$ .
  - ▶ equivalently, let  $w_i$  be a one-hot vector (dimensionality  $n_w = \text{vocab size}$ ) where the associate word's index equals one .

# Word Embeddings

- ▶ Embedding layer maps word indexes to dense vectors.
- ▶ Documents are lists of word indexes  $\{w_1, w_2, \dots, w_{n_i}\}$ .
  - ▶ equivalently, let  $w_i$  be a one-hot vector (dimensionality  $n_w = \text{vocab size}$ ) where the associate word's index equals one .
  - ▶ Normalize all documents to the same length  $L$ ; shorter documents can be padded with a null token. This requirement can be relaxed with recurrent neural networks.

# Word Embeddings

- ▶ Embedding layer maps word indexes to dense vectors.
- ▶ Documents are lists of word indexes  $\{w_1, w_2, \dots, w_{n_i}\}$ .
  - ▶ equivalently, let  $w_i$  be a one-hot vector (dimensionality  $n_w = \text{vocab size}$ ) where the associate word's index equals one .
  - ▶ Normalize all documents to the same length  $L$ ; shorter documents can be padded with a null token. This requirement can be relaxed with recurrent neural networks.
- ▶ The embedding layer replaces the list of sparse one-hot vectors with a list of  $n_E$ -dimensional ( $n_E \ll n_w$ ) dense vectors

$$\mathbf{X} = \begin{bmatrix} x_1 & \dots & x_L \end{bmatrix}$$

where

$$\underbrace{x_i}_{n_E \times 1} = \underbrace{\Omega'}_{n_E \times n_w} \cdot \underbrace{w_i}_{n_w \times 1}$$

# Word Embeddings

- ▶ Embedding layer maps word indexes to dense vectors.
- ▶ Documents are lists of word indexes  $\{w_1, w_2, \dots, w_{n_i}\}$ .
  - ▶ equivalently, let  $w_i$  be a one-hot vector (dimensionality  $n_w = \text{vocab size}$ ) where the associate word's index equals one .
  - ▶ Normalize all documents to the same length  $L$ ; shorter documents can be padded with a null token. This requirement can be relaxed with recurrent neural networks.
- ▶ The embedding layer replaces the list of sparse one-hot vectors with a list of  $n_E$ -dimensional ( $n_E \ll n_w$ ) dense vectors


$$\mathbf{X} = \begin{bmatrix} x_1 & \dots & x_L \end{bmatrix}$$

where

$$\underbrace{x_i}_{n_E \times 1} = \underbrace{\Omega'}_{n_E \times n_w} \cdot \underbrace{w_i}_{n_w \times 1}$$

- ▶ This  $\mathbf{X}$  matrix is then flattened into a  $L * n_E$  vector for input to the next layer.

# Illustration



# Examining the Embeddings

- ▶ See Jupyter notebook for examples on training and visualizing the embeddings with words as points.
  - ▶ Also examples for extracting vectors for words and computing cosine similarity between words.

# Word Embeddings – Word2Vec, GloVe

- ▶ Word embeddings:
  - ▶ refers to a class of statistical models that **represent words or phrases as points in a vector space**.

# Word Embeddings – Word2Vec, GloVe

- ▶ Word embeddings:
  - ▶ refers to a class of statistical models that **represent words or phrases as points in a vector space**.
- ▶ The key idea is to represent the meaning of words by the neighboring words – their **contexts**.



# Word Embeddings – Word2Vec, GloVe

- ▶ Word embeddings:
  - ▶ refers to a class of statistical models that **represent words or phrases as points in a vector space**.
- ▶ The key idea is to represent the meaning of words by the neighboring words – their **contexts**.
- ▶ You might hear “word embeddings” and “word2vec” interchangeably, although word2vec technically refers to a particular implementation of a word embedding model.
  - ▶ the other well-known implementation is gloVe, which is faster but has similar performance/applications

# Word Embeddings and Word2Vec

- ▶ Word2Vec , GloVe, and other popular embeddings vectors are trained the same way as the word embeddings we just made for citation counts.
  - ▶ rather than predicting some metadata (such as citations) they predict the co-occurrence of neighboring words.

# Why word vectors?

- ▶ Once words are represented as vectors, we can use linear algebra to understand the relationships between words:
  - ▶ Words that are geometrically close to each other are similar: e.g. “student” and “pupil.”

## Why word vectors?

- ▶ Once words are represented as vectors, we can use linear algebra to understand the relationships between words:
  - ▶ Words that are geometrically close to each other are similar: e.g. “student” and “pupil.”
  - ▶ More intriguingly, word2vec algebra can depict conceptual, analogical relationships between words.
    - ▶ Consider the analogy: **man is to king as woman is to \_\_\_\_\_**

# Why word vectors?

- ▶ Once words are represented as vectors, we can use linear algebra to understand the relationships between words:
  - ▶ Words that are geometrically close to each other are similar: e.g. “student” and “pupil.”
  - ▶ More intriguingly, word2vec algebra can depict conceptual, analogical relationships between words.
    - ▶ Consider the analogy: **man is to king as woman is to \_\_\_\_\_**
    - ▶ With word2vec, we have

$$\text{vec}(\text{king}) - \text{vec}(\text{man}) + \text{vec}(\text{woman}) \approx \text{vec}(\text{queen})$$

## How are word embeddings different from topic models?

- ▶ Ben Schmidt:
  - ▶ Topic models reduce words to core meanings to understand documents more clearly.

# How are word embeddings different from topic models?

- ▶ Ben Schmidt:
  - ▶ Topic models reduce words to core meanings to understand documents more clearly.
  - ▶ Word embedding models ignore information about individual documents to better understand the relationships between words.

## Word Function $\longleftrightarrow$ Word Neighbors

- ▶ "The meaning of a word is its use in the language"

- Ludwig Wittgenstein, *Philosophical Investigation*, 1953



## Word Function $\longleftrightarrow$ Word Neighbors

- ▶ "The meaning of a word is its use in the language"

- Ludwig Wittgenstein, *Philosophical Investigation*, 1953


- ▶ "You shall know a word by the company it keeps"

- J.R. Firth, *Papers in Linguistics*, 1957

I've never seen this word before, but...

- ▶ He filled the **wampimuk**, passed it around and we all drunk some
- ▶ We found a little, hairy **wampimuk** sleeping behind the tree


# Words as Vectors




- Use cosine similarity as a measure of relatedness:

$$\cos \theta = \frac{v_1 \cdot v_2}{||v_1|| ||v_2||}$$


# Vector Directions $\leftrightarrow$ Meaning



Male-Female



Verb Tense



Country-Capital

# Linguistic Relations


## Synonymy



## Antonymy



## Hyponymy



## Collocational Relations

### Collocation

<i>against the</i>	<b><i>law</i></b>	
	<b><i>law</i></b>	<i>enforcement</i>
<i>become</i>	<b><i>law</i></b>	
	<b><i>law</i></b>	<i>is passed</i>

### Colligation

<i>normal</i>	<i>VERB past</i>	<b><i>time</i></b>
	<i>saved</i>	
	<i>spent</i>	
	<i>wasted</i>	
<i>sport</i>	<i>ADJECTIVE</i>	<b><i>time</i></b>
	<i>half</i>	
	<i>extra</i>	
	<i>full</i>	

# Similarity vs. Relatedness

- ▶ Semantic **similarity**: words sharing salient attributes / features
  - ▶ synonymy (car / automobile)
  - ▶ hypernymy (car / vehicle)
  - ▶ co-hyponymy (car / van / truck)

# Similarity vs. Relatedness

- ▶ Semantic **similarity**: words sharing salient attributes / features
  - ▶ synonymy (car / automobile)
  - ▶ hypernymy (car / vehicle)
  - ▶ co-hyponymy (car / van / truck)
- ▶ Semantic **relatedness**: words semantically associated without necessarily being similar
  - ▶ function (car / drive)
  - ▶ meronymy (car / tire)
  - ▶ location (car / road)
  - ▶ attribute (car / fast)

(Budansky and Hirst, 2006)



## Most similar words to dog, depending on window size

	2-word window	30-word window	
More paradigmatic		<u>kennel</u>	More syntagmatic
	cat	puppy	
	horse	pet	
	fox	bitch	
	pet	terrier	
	rabbit	rottweiler	
	pig	canine	
	animal	cat	
	mongrel	<u>bark</u>	
	sheep	alsatian	
	pigeon		

- ▶ Small windows pick up substitutable words; large windows pick up topics.

# Evaluation of Word Embeddings

- ▶ Intrinsic:
  - ▶ evaluate word-pairs similarities → compare with similarity judgments given by humans
  - ▶ evaluate on analogy tasks (“Paris is to France as Tokyo is to \_\_\_\_”)

# Evaluation of Word Embeddings

- ▶ Intrinsic:
  - ▶ evaluate word-pairs similarities → compare with similarity judgments given by humans
  - ▶ evaluate on analogy tasks (“Paris is to France as Tokyo is to \_\_\_\_”)
- ▶ Extrinsic:
  - ▶ use the vectors in a downstream task (classification, translation, ...) and evaluate the final performance on the task

## SGNS: Skip-gram with negative sampling

- ▶ When people mention “word2vec”, they are usually talking about SGNS: “skip gram with negative sampling.”
  - ▶ This is a particular word-embedding model with good performance on a range of analogy and prediction tasks.

## SGNS: Skip-gram with negative sampling

- ▶ When people mention “word2vec”, they are usually talking about SGNS: “skip gram with negative sampling.”
  - ▶ This is a particular word-embedding model with good performance on a range of analogy and prediction tasks.
- ▶ How does it learn the meaning of the word “fox”:

The quick brown **fox** jumps over the lazy dog

## SGNS: Skip-gram with negative sampling


- ▶ When people mention “word2vec”, they are usually talking about SGNS: “skip gram with negative sampling.”
  - ▶ This is a particular word-embedding model with good performance on a range of analogy and prediction tasks.

- ▶ How does it learn the meaning of the word “fox”:

The quick brown **fox** jumps over the lazy dog

- ▶ Word2Vec reads in every example of the word “fox” and tries to predict what other words will be in the context window.
  - ▶ the prediction weights on these other words (after dimension reduction) are the word vectors

# Word2Vec Schema



# Tokenizing for Embeddings

- ▶ embeddings work better with more information about the placement of words in sentences.
  - ▶ don't drop stopwords/function-words
  - ▶ should include tokens for start of sentence and end of sentence
  - ▶ should include a special token for out-of-vocabulary words
    - ▶ or replace with the part of speech tag



# Word Dropout

- ▶ When training models, words can be randomly replaced with the unknown symbol with some small probability (Iyyer et al 2015).
- ▶ Prevents models from relying too much on particular words.

# K-means clustering with Word Embeddings

## Income Tax (Pensions Topic and Health Care Topic)

pension\_board have\_attain\_the\_age  
in\_excess\_of\_year  
retir\_purpos  
such\_depend  
such\_servic  
biweek\_pay\_period

medicar servic in accord vocat rehabilit  
legal\_settlement  
admiss\_center  
self-support  
depend\_children  
coron\_medic\_condit  
babiday\_servic  
cerebr\_palsi

## Sales Tax (Retail Topic and Health Care Topic)

fuel\_dealer  
retail\_store  
such\_distributor  
har\_race

psychiatrist juvenil\_offend  
state\_plan  
educ\_or\_train  
retard\_servic  
aid\_to\_famili  
cost\_of\_health  
first\_aid

- Clustered phrases affecting tax revenues (Ash 2018); Green words tend to increase revenues; red words tend to decrease revenues.

## Word Mover Distance

- ▶ Cosine distance treats synonyms as just as close as totally unrelated words.

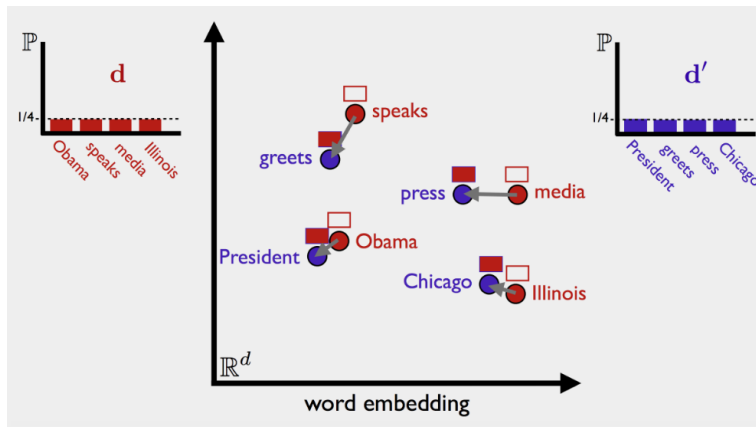
# Word Mover Distance

- ▶ Cosine distance treats synonyms as just as close as totally unrelated words.
- ▶ Word mover distance between two texts is given by:
  - ▶ total amount of “mass” needed to move words from one side into the other
  - ▶ multiplied by the distance the words need to move
  - ▶ Kusner, Sun, Kolkin, and Weinberger (ICML 2015)

# Word Mover Distance

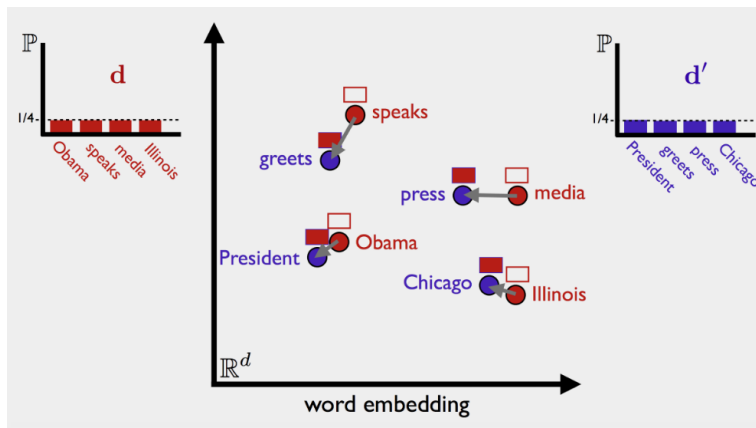
- ▶ Cosine distance treats synonyms as just as close as totally unrelated words.
- ▶ Word mover distance between two texts is given by:
  - ▶ total amount of “mass” needed to move words from one side into the other
  - ▶ multiplied by the distance the words need to move
  - ▶ Kusner, Sun, Kolkin, and Weinberger (ICML 2015)
- ▶ Requires measure of distance between words (word embeddings).
  - ▶ see wmd package in Python.

# Illustration



- ▶  $d$  (obama speaks media illinois) is orthogonal to  $d'$  (president greet press chicago):
  - ▶ cosine similarity is zero

# Illustration



- ▶  $d$  (obama speaks media illinois) is orthogonal to  $d'$  (president greets press chicago):
  - ▶ cosine similarity is zero
  - ▶ Word mover distance sums the shortest distances between the words in the documents.

## Pre-trained word embeddings

- ▶ For some NLP tasks, you might not need to train your own vectors.
  - ▶ or your corpus might be too small to do so.



## Pre-trained word embeddings

- ▶ For some NLP tasks, you might not need to train your own vectors.
  - ▶ or your corpus might be too small to do so.
- ▶ In these cases, you can use a pre-trained model, like spaCy's GloVe embeddings:
  - ▶ one million vocabulary entries
  - ▶ 300-dimensional vectors
  - ▶ trained on the Common Crawl corpus

## Pre-trained word embeddings

- ▶ For some NLP tasks, you might not need to train your own vectors.
  - ▶ or your corpus might be too small to do so.
- ▶ In these cases, you can use a pre-trained model, like spaCy's GloVe embeddings:
  - ▶ one million vocabulary entries
  - ▶ 300-dimensional vectors
  - ▶ trained on the Common Crawl corpus
- ▶ Can initialize prediction model using pre-trained embeddings.

## Tips for using pre-trained embeddings

- ▶ Split training in two steps:
  - ▶ in first run, train the model with the first layer (the pre-trained embeddings) frozen.

## Tips for using pre-trained embeddings

- ▶ Split training in two steps:
  - ▶ in first run, train the model with the first layer (the pre-trained embeddings) frozen.
  - ▶ in second run, un-freeze the embedding layer for fine tuning.

# Outline

Intro to Neural Nets

Practicalities

Autoencoders

Embedding Layers

Word Embeddings

Kozlowski, Evans, and Taddy (2019)

## Kozlowski, Evans, and Taddy (2019)

- ▶ Data-set: Google 5-grams.
  - ▶ n-grams of length five for U.S. and U.K. publications.
  - ▶ provides counts for each year

## Kozlowski, Evans, and Taddy (2019)

- ▶ Data-set: Google 5-grams.
  - ▶ n-grams of length five for U.S. and U.K. publications.
  - ▶ provides counts for each year
- ▶ Extract language dimensions:
  - ▶ get the complete list of WordNet antonym pairs (e.g, “weak/strong”, “tall/short”)
    - ▶ filter on document frequency to 428 pairs.
  - ▶ map the dimensional shifts between the antonyms.

## Kozlowski, Evans, and Taddy (2019)

- ▶ Data-set: Google 5-grams.
  - ▶ n-grams of length five for U.S. and U.K. publications.
  - ▶ provides counts for each year
- ▶ Extract language dimensions:
  - ▶ get the complete list of WordNet antonym pairs (e.g, “weak/strong”, “tall/short”)
    - ▶ filter on document frequency to 428 pairs.
  - ▶ map the dimensional shifts between the antonyms.
  - ▶ compare this vector shift to the one between men and women.



## Mapping gender, class, and race

Gender	Class	Race <sup>†</sup>
man – woman	rich – poor	black – white
men – women	richer – poorer	blacks – whites
he – she	richest – poorest	Blacks – Whites
him – her	affluence – poverty	Black – White
his – her	affluent – impoverished	African – European
his – hers	expensive – inexpensive	African – Caucasian
boy – girl	luxury – cheap	
boys – girls	opulent – needy	
male – female		
masculine – feminine		

## Matching antonyms to gender/class

### Gender dimension nearest neighbors

- |                    |                         |
|--------------------|-------------------------|
| 1. rugged–delicate | .219<br>(.213, .224)    |
| 2. soft–loud       | -.209<br>(-.216, -.201) |
| 3. tender–tough    | -.202<br>(-.210, -.197) |
| 4. timid–bold      | -.181<br>(-.186, -.174) |
| 5. soft–hard       | -.161<br>(-.168, -.158) |

### Class dimension nearest neighbors

- |                          |                         |
|--------------------------|-------------------------|
| 1. weak-strong           | -.292<br>(-.301, -.287) |
| 2. fortunate-unfortunate | .291<br>(.286, .297)    |
| 3. unhappy-happy         | -.259<br>(-.266, -.254) |
| 4. beautiful-ugly        | .242<br>(.238, .245)    |
| 5. potent_impotent       | .234<br>(.227, .244)    |

# Mapping musical genres to race/class

